# LinkedIn Profile Audit Report

Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics

Generated August 28, 2026

OVERVIEW

## Executive Summary

Felipe's profile tells a coherent, increasingly rare story: an analyst who moved from BI reporting (Renoboy, early Blackboard) into governed, high-scale Microsoft Fabric engineering (EPAM), backed by a strong paired DP-700/DP-600 certification set, heavily quantified achievements ($45M in opportunity generation, $10M in campaign-driven revenue, 35% latency reductions), specific and credible received recommendations, and an active, on-brand posting cadence (one post reached 199 reactions and 44K+ impressions). Both career-coach and recruiter agree the underlying narrative and screenability are strong, but four issues are flagged independently by both agents and therefore carry outsized risk: the headline title doesn't exactly match the formal EPAM job title, the About section ends with no call to action, the EPAM and Blackboard Senior Data Analyst entries are dense, un-trimmed blocks that bury strong metrics from a recruiter's six-second skim, and Jason Bass's recommendation references an 'Anthology' employer that never appears anywhere else in the profile. None of these undermine the strong Fabric-engineer positioning, but each costs clarity or trust at first read. Recommended path forward: stay an individual contributor and deepen into AI/GenAI engineering on top of the existing Fabric base, while retaining the Fabric/Medallion depth that currently differentiates the profile.

SNAPSHOT

## Profile Snapshot

- **Header:** present
- **About:** present
- **Top Skills:** 5 captured
- **Experience:** 4 captured
- **Education:** 2 captured
- **Skills Full List:** 24 captured
- **Certifications:** 3 captured
- **Featured Items:** 2 captured
- **Recent Posts:** 3 captured
- **Recommendations Received:** 3 captured
- **Recommendations Given:** 2 captured
- **Projects:** 4 captured
- **Missing:** none — all sections captured

FINDINGS

## Section-by-Section Findings

### About

- **CAREER COACH** — The opening line ('8+ years building the data infrastructure that powers strategic decisions') is serviceable but leans on familiar LinkedIn phrasing ('powers strategic decisions') rather than a distinct, specific claim. It states years of experience and a vague value claim before earning the reader's attention with a concrete detail.

- **CAREER COACH** — The About section ends on a certifications list with no call to action — no invitation to connect, no statement of what kind of role/conversation Felipe is open to. For a profile clearly built to attract inbound recruiter or hiring-manager interest, this is a real gap, not just a stylistic nice-to-have.
- **RECRUITER** — The about section is well-written and keyword-rich but ends on the DP-700/DP-600 certification sentence with no call-to-action (e.g., no invitation to connect, no statement of what Felipe is looking for next, no contact prompt). A recruiter skimming to the end gets no explicit next step.

## Positioning & Narrative

- **CAREER COACH** — The headline reads 'Senior Analytics Engineer' but the current EPAM title in the experience section is 'Senior Data Analytics Engineer.' It's a minor wording gap, not a level or domain mismatch, but a recruiter scanning both fields side by side will notice the titles don't match exactly.
- **CAREER COACH** — The headline and About section both commit hard to a specific, well-defined niche — Microsoft Fabric, Medallion Lakehouse architecture, semantic layer, real-time analytics — rather than reading as a generic 'data professional' profile. Seniority (Senior) and domain (Fabric-centric data/analytics engineering) are unambiguous.

## Experience

- **CAREER COACH** — Bullets consistently quantify outcomes rather than just listing tasks: 35% latency reduction, 40% faster incident response, 35% faster report load times, $45M in opportunity generation, $10M in campaign revenue, 80% cross-sell accuracy. This is strong, results-oriented framing throughout the experience section.
- **CAREER COACH** — The EPAM and both Blackboard entries pack 6+ dense, long-sentence bullets plus intro/outro paragraphs into single description blocks. Individual bullets like the EPAM 'Delivered 20+ enterprise reports...' line run 40+ words with multiple clauses, which undermines scannability even though the underlying content is strong.

## Skills & Keywords

- **CAREER COACH** — top_skills surfaces only 5 items (SQL, Python, Microsoft Fabric, Power BI, Azure), all consistent with the stated niche and not padded with irrelevant keywords. The fuller skills_full_list is Fabric/Microsoft-ecosystem-only, with no competing platform (e.g., Databricks, Snowflake) represented anywhere — a positioning choice that reinforces focus but narrows perceived platform range.
- **RECRUITER** — header.headline ('Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics') is keyword-rich and aligns well with top_skills and the terminology used throughout about and experience (Medallion, semantic layer, Power BI, Fabric, DAX, real-time). This is a strength for recruiter keyword searches on those terms.
- **RECRUITER** — header.headline states the title 'Senior Analytics Engineer', but the current role in experience is formally titled 'Senior Data Analytics Engineer' at EPAM Systems. A recruiter or ATS doing an exact-title keyword search for 'Data Analytics Engineer' may not surface this profile as a strong title match, and a human reviewer comparing the two fields side by side would notice the discrepancy within seconds.
- **RECRUITER** — top_skills (SQL, Python, Microsoft Fabric, Microsoft Power BI, Microsoft Azure) are a coherent subset of skills_full_list and are echoed throughout about and experience prose (PySpark, Delta Tables, DAX, semantic models) — no evidence of keyword-stuffing or unsupported claims.
- **RECRUITER** — skills_full_list (24 entries) is tightly curated around the Fabric/Power BI/data engineering stack (Medallion Architecture, Delta Lake, PySpark, Apache Spark, DAX, Data Governance, Data Modeling, ETL, Azure Data Factory, etc.). No outdated or irrelevant entries dilute the signal, and every top_skills entry is also present in skills_full_list.

## Certifications

- **CAREER COACH** — All three certifications (DP-700, DP-600, Power BI Data Analyst) are current Microsoft credentials, all valid through 2027, and map precisely onto the profile's stated Fabric engineering + analytics positioning. This is a coherent, well-maintained credential set rather than a scattershot list.

## Featured

- **CAREER COACH** — Recent posts on technical topics (Fabric Warehouse MERGE dedup, PySpark execution plans, memory-constrained CSV loads) are getting real engagement, including one post at 199 reactions and 44K+ impressions, which reinforces the hands-on engineering credibility the rest of the profile claims. Only 2 featured items are pinned despite 3+ posts and 4 projects existing, so the featured section is under-curated relative to available material.

## Projects

- **CAREER COACH** — The two most recent projects (NYC Taxi Medallion, Real-Time Stock Volatility) directly extend the Fabric/streaming positioning with fresh, dated work (through Mar 2026). However, the Market Basket Analysis project explicitly mirrors an achievement already claimed at Blackboard ('This mirrors the methodology behind... at Blackboard'), which risks reading as recycled credit rather than new independent depth to a skeptical reader.

## Recommendations

- **CAREER COACH** — Jason Bass's recommendation states Felipe 'worked on my team for several years at Blackboard and then Anthology,' but 'Anthology' never appears anywhere in the experience, about, or projects sections — both Blackboard roles (Nov 2019–Mar 2024) are listed under 'Blackboard' only. This leaves an unreconciled company-naming gap that a reader unfamiliar with the real-world Blackboard-to-Anthology rebrand will find confusing.
- **CAREER COACH** — Both recommendations given (to Uladzimir Kruhlik at EPAM, to Diego Sanabria at Blackboard) are internally consistent with the experience timeline — the Uladzimir recommendation's DP-600 mention aligns with the Sep 2024 certification date during the EPAM tenure, and Diego's mentorship reference to the 2020 pandemic-era $45M outcome falls correctly within the Nov 2019–Jan 2022 Data Analyst role. These reinforce the mentorship and continuous-learning thread running through the profile.

## Posts & Visibility

- **RECRUITER** — recent_posts contains 3 entries dated 6d, 1w, and 2w ago — a healthy, active posting cadence that signals recent engagement with the platform rather than a stale/dormant profile.
- **RECRUITER** — All three recent_posts are technical deep-dives on Fabric, PySpark, and data engineering troubleshooting (execution plans, MERGE/dedupe logic, memory-constrained CSV processing) — directly reinforcing the headline's positioning around Microsoft Fabric and real-time/engineering-heavy analytics work. Topics are on-brand and would read as credible thought leadership to a recruiter sourcing for this profile.
- **RECRUITER** — Engagement is inconsistent across the three posts: the most recent post (6d ago) has only 2 likes, 0 comments, and 1,147 impressions, versus 177 likes/47,595 impressions (1w) and 199 likes/44,025 impressions (2w) for the prior two. The steep drop-off on the newest post is a weaker credibility/engagement signal relative to the other two, though not necessarily indicative of a broader trend on a 3-post sample.

## Credibility & Red Flags

- **RECRUITER** — Three Microsoft certifications are listed (Fabric Data Engineer Associate, Fabric Analytics Engineer Associate, Power BI Data Analyst Associate), all directly relevant to the target role, with clear issue and expiration dates, and all currently unexpired (latest expiration Nov 2027 / Sep 2027).

- **RECRUITER** — Achievements across all four roles are heavily quantified (e.g., '35% improvement in report load times', '40% reduction in response time', '$45M in opportunity generation', '$10M in campaign-driven revenue', '80% accuracy', '20+ enterprise reports', 'billions of records', 'team of 3 Data Engineers'), which reads as strong, skimmable evidence of impact rather than vague self-description.
- **RECRUITER** — All three recommendations_received (Claudio Warneboldt, Jason Bass, Andy McGinn) are specific and detailed, citing concrete projects (e.g., the GTM Power BI report, the analytical foundation for the Open LMS/Blackboard Collaborate divestiture) rather than generic praise — a genuine credibility signal.
- **RECRUITER** — Jason Bass's recommendation states Felipe 'worked on my team for several years at Blackboard and then Anthology,' but 'Anthology' does not appear anywhere in header, about, or experience — both Blackboard roles (Data Analyst and Senior Data Analyst) are listed only as 'Blackboard' with no mention of a name change or a separate Anthology entry. A recruiter cross-checking the recommendation against the experience section would notice this gap and may question whether an employer/timeline is missing or unexplained.

### Six Second Skim

- **RECRUITER** — header.photo_present and header.banner_present are both true, and the headline is specific rather than a generic title — no red flag on visual/first-impression completeness.
- **RECRUITER** — Experience timeline is fully continuous with no unexplained gaps: Renoboy S.A (Dec 2017–Nov 2019) → Blackboard Data Analyst (Nov 2019–Jan 2022) → Blackboard Senior Data Analyst (Jan 2022–Mar 2024) → EPAM Senior Data Analytics Engineer (Apr 2024–Present).
- **RECRUITER** — The EPAM Systems and Blackboard 'Senior Data Analyst' entries are each structured as an opening narrative paragraph, 5-6 bullet points, and a closing narrative paragraph — a dense, un-trimmed wall-of-text format compared to the more skimmable bullet-forward Blackboard 'Data Analyst' and Renoboy entries. A recruiter skimming for 6 seconds is more likely to read only the first line and miss the strongest quantified bullets buried mid-block.

PRIORITIES

## Prioritized Recommendations

### High Impact

#### Headline / Current Title Mismatch (High Impact)

Both agents independently flag that the headline title 'Senior Analytics Engineer' doesn't exactly match the formal EPAM title 'Senior Data Analytics Engineer.' Career-coach frames it as a narrative consistency gap; recruiter frames the same gap as an exact-title keyword/ATS search risk that a human reviewer would also notice within seconds. Confirmed by both agents from two independent angles, bumping this to high impact.

**Before**

Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics

**After**

Senior Data Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics

#### About Section Missing a Call to Action (High Impact)

Career-coach rates this high severity as a real gap for a profile built to attract inbound interest — the About section ends on a certifications list with no invitation to connect. Recruiter independently flags the same gap from a six-second-skim standpoint, noting a recruiter reaches the end of About with no explicit next step. Confirmed by both agents.

**Before**

I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack.

**After**

I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack. I'm always open to connecting with teams building AI-ready data platforms on Microsoft Fabric — message me if you're tackling similar Medallion or real-time analytics challenges.

### Dense, Un-trimmed EPAM and Blackboard Experience Bullets (High Impact)

Recruiter rates this a high-severity six-second-skim failure: the EPAM and Blackboard Senior Data Analyst entries are structured as narrative-intro-plus-bullets-plus-narrative-outro blocks that bury the strongest quantified metrics. Career-coach independently flags the same density/readability issue at medium severity, citing 40+ word, multi-clause bullets such as the EPAM report-delivery line. Confirmed by both agents, bumped to high impact.

**Before**

Delivered 20+ enterprise reports serving Finance, FP&A, Marketing, Strategy, and Sales teams by optimizing semantic models and Power BI performance using advanced DAX and UI/UX best practices, achieving a 35% improvement in report load times for executive stakeholders.

**After**

Delivered 20+ enterprise Power BI reports for Finance, FP&A, Marketing, Strategy, and Sales. Improved report load times 35% via semantic model and DAX optimization.

### Unreconciled "Anthology" Reference in Jason Bass's Recommendation (High Impact)

Both agents flag that Jason Bass's recommendation states Felipe worked at 'Blackboard and then Anthology,' but 'Anthology' never appears anywhere in the header, about, experience, or projects sections — both Blackboard roles are listed only as 'Blackboard.' This leaves the real-world Blackboard-to-Anthology rebrand unreconciled and forces a reader to guess at an unexplained employer name. Confirmed independently by both agents.

**Before**

Blackboard

**After**

Blackboard (rebranded Anthology in 2021)

## Low Impact

### About Section Opening Hook Leans on Generic Phrasing (Low Impact)

Career-coach (low severity) notes the About section's opening line leans on familiar LinkedIn phrasing ('powers strategic decisions') rather than leading with a specific, differentiating fact, stating years of experience and a vague value claim before earning the reader's attention. Not flagged by recruiter.

**Before**

8+ years building the data infrastructure that powers strategic decisions.

**After**

I hold both Microsoft Fabric certifications (DP-700 and DP-600) and have built Medallion Lakehouses processing billions of records for Finance and FP&A teams at EPAM.

### Market Basket Analysis Project Reads as Recycled Credit (Low Impact)

Career-coach (low severity) notes the Market Basket Analysis project description explicitly states it 'mirrors the methodology behind... at Blackboard,' which risks reading as recycled credit rather than new independent depth to a skeptical reader. The suggested fix is a reordering of the description (lead with what's new before referencing Blackboard) rather than a specific rewritten passage, so no concrete before/after text is available without inventing content. Not flagged by recruiter.

### Featured Section Under-curated Relative to Available Material (Low Impact)

Career-coach (low severity) notes only 2 items are pinned to Featured despite 3+ recent posts (including one at 199 reactions and 44K+ impressions) and 4 projects existing, leaving the profile's strongest proof points requiring extra scrolling to find. This is a curation/selection action rather than a text edit, so no before/after text applies. Not flagged by recruiter.

STRATEGY

## Career Trajectory — 3-Year Outlook

Forward-looking strategic guidance on the candidate's career trajectory — distinct from the profile-polish recommendations above.

### SWOT Analysis

#### Strengths

- Holds both DP-700 and DP-600, covering Fabric's engineering and analytics layers, a rare paired credential combination.
- Led a 3-person engineering team while shipping 100+ production Delta tables and real-time pipelines that cut latency 35%.
- Independent project portfolio spans Fabric engineering (Eventstream, PySpark) and applied ML (XGBoost, Apriori), showing range beyond pure BI reporting.

#### Opportunities

- Roughly two-thirds of enterprises now run multiple data platforms by design, so pairing Fabric depth with Databricks exposure fits real hybrid-stack hiring.
- AI Engineer job postings are growing far faster than analytics roles, and his EPAM work already frames Fabric as groundwork for Generative AI.
- Data engineer salaries are trending toward $153K in 2026, and DP-700 is the exact certification Microsoft is currently steering hiring demand toward.

#### Weaknesses

- No Databricks, Snowflake, or non-Microsoft lakehouse tool appears anywhere in the skills list, leaving platform exposure narrow.
- No LangChain, vector database, or LLM-orchestration skill is listed despite the About section referencing future Generative AI work.
- Team leadership appears only once, a 3-person team at EPAM, with no formal management title or certification behind it.

#### Threats

- Reports on Azure's job-posting share conflict, one citing a drop from 75% to 34% and another citing 74.5% demand, so treat this signal as weak.
- Databricks holds a stronger enterprise reputation rating than Fabric in analyst reviews, a signal complex engineering workloads may keep favoring it.
- Analytics Engineer pay trails AI Engineer pay by roughly $50K, so staying purely in BI-adjacent analytics engineering risks a widening compensation gap.

## Recommended Paths

### Stay individual contributor and deepen into GenAI/AI engineering on top of the existing Fabric base.

**Rationale:**

My recommendation: AI Engineer demand and pay are growing far faster than analytics engineering, and his EPAM description already frames Fabric as the 'foundational layer for future Generative AI and LLM integrations' — the churn-prediction project shows he's already reaching past pure BI into ML.

**Tradeoffs:**

Requires building real LangChain/RAG project evidence from zero, since none exists in the current skill or project list, and moving too fast could dilute the sharp Fabric-engineer brand that's currently working well.

### Add Databricks depth alongside Fabric to become a platform-agnostic lakehouse engineer.

**Rationale:**

My recommendation: enterprises increasingly run Fabric and Databricks side by side by design, and Databricks carries a stronger engineering reputation in analyst reviews, so adding it would reduce single-vendor risk in an otherwise Microsoft-only skill set.

**Tradeoffs:**

Splits certification and project time across two ecosystems instead of compounding depth in one, and EPAM's current Fabric-centric client work may not provide hands-on Databricks exposure to back it up.

### Pursue a Data/Analytics Engineering Manager track building on the existing team-lead experience.

**Rationale:**

My recommendation: he already led a 3-person team at EPAM and has a consistent mentoring pattern across Blackboard and EPAM (per both recommendations given), which is real, if thin, management signal to build on.

**Tradeoffs:**

Management-track compensation growth is generally slower than the AI-specialist pay growth seen in search results, and shifting focus toward people management would reduce the hands-on technical depth that currently differentiates his profile.

## Skill & Certification Recommendations

### LangChain/RAG fundamentals plus one small agentic or retrieval project (High Priority)

It directly extends the Generative AI groundwork already claimed in his About/EPAM description with real evidence, in the fastest-growing demand segment found in search.

### Databricks Certified Data Engineer Associate (Medium Priority)

Closes the one major platform gap in an otherwise all-Microsoft skill list and matches the hybrid-platform hiring pattern enterprises are converging on.

### AI governance / responsible AI practices (Medium Priority)

He already built a company-wide Data Governance framework at EPAM, so extending that into AI governance is a natural adjacency rather than a generic buzzword chase.

### A second/duplicate Power BI or DAX credential (Low Priority)

Not worth pursuing — he already holds PL-300 plus two Fabric certifications that thoroughly cover the BI/semantic layer.

### Formal people-management training or a technical-lead certificate (Low Priority)

Only worth prioritizing if he commits to the management path above; his current leadership evidence is a single 3-person team, too thin to invest here yet.

### OODA Execution Loop

**Observe**

Each quarter, review which of his own project types (Fabric/streaming engineering vs. Python/ML projects like the churn and market-basket work) drew more recruiter messages, post engagement, or interview requests, and track DP-700/DP-600 renewal dates (Nov 2027, Sep 2027).

→

**Orient**

Compare that engagement signal against the three recommended paths above — is the market pulling him toward GenAI/AI engineering, platform-agnostic data engineering, or a leadership track — and re-run a light market search (AI Engineer, Data Engineer, Analytics Manager demand) to check the SWOT findings still hold.

→

**Decide**

Pick one concrete direction for the next quarter's effort rather than splitting time across all three paths at once, based on whichever showed the strongest pull in Orient; revisit the decision every quarter rather than locking in for the full 3 years upfront.

→

**Act**

Ship one tangible artifact each quarter tied to the chosen direction — a LangChain/RAG project, a Databricks-based project, or a documented leadership outcome — and update the headline/About copy to reflect the current direction so the profile never runs three unresolved signals in parallel.

ROADMAP

## Next Steps

### This Week

- Align the headline title with EPAM's formal title, 'Senior Data Analytics Engineer' (Headline / Current Title Mismatch recommendation).
- Add a closing call-to-action sentence to the end of the About section (About Section Missing a Call to Action recommendation).
- Add a '(rebranded Anthology in 2021)' parenthetical to both Blackboard experience entries to reconcile Jason Bass's recommendation text (Unreconciled 'Anthology' Reference recommendation).
- Rewrite the About section's opening line to lead with a specific differentiator instead of generic framing (About Section Opening Hook recommendation).

### This Month

- Restructure the EPAM and Blackboard Senior Data Analyst experience entries: trim the narrative intro/outro paragraphs and split dense, multi-clause bullets into single-idea lines (Dense, Un-trimmed EPAM and Blackboard Experience Bullets recommendation).
- Pin 1-2 additional items to the Featured section — e.g., the highest-engagement post and one of the two 2026 Fabric projects — so top proof points are visible without scrolling (Featured Section Under-curated recommendation).
- Reorder the Market Basket Analysis project description to lead with what's genuinely new in the personal version before referencing the Blackboard connection (Market Basket Analysis Project Reads as Recycled Credit recommendation).

### This Quarter

- As new posts or projects are published each quarter, keep revisiting Featured curation so it continues to reflect the two strongest current proof points (Featured Section Under-curated recommendation).
- As additional independent projects are added over time, continue auditing project descriptions for recycled-credit framing so each reads as new depth rather than a restated employer achievement (Market Basket Analysis Project Reads as Recycled Credit recommendation).