LinkedIn Profile Audit Report

Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics

Generated August 30, 2026

OVERVIEW

Executive Summary

Felipe's profile presents a coherent, well-differentiated Senior Analytics Engineer story: headline, About, experience, skills, and certifications consistently reinforce a Microsoft Fabric / Power BI / Medallion-architecture specialization, backed by strong quantified achievements in every role (35% latency reduction, 55% faster pipeline delivery, $45M in opportunity generation, $210M and $31M divestiture support, 38% churn reduction, 40% forecast accuracy improvement), three current and directly relevant Microsoft certifications, three substantive named recommendations, and an active, on-topic posting cadence including two standout posts (177 and 199 likes, ~44K-48K impressions each). Both the career-coach and recruiter lenses independently rate this as a strong profile out of the gate for narrative coherence and 6-second screenability. The headline issues both perspectives converge on are: an unaddressed pivot into multi-agent/Claude Code AI-tooling work (visible in two of the newest projects and two skills entries but never acknowledged in the headline or About, which simultaneously dilutes recruiter keyword focus and leaves the most current work unexplained to any reader), an unreconciled company-name discrepancy between a received recommendation (which mentions working together at 'Anthology') and the experience section (which lists only 'Blackboard'), and a uniformly templated intro-bullets-closing structure repeated across all four roles that reads as formulaic and adds scan friction for a recruiter skim. Secondary gaps include a top_skills list that omits heavily-used terms like PySpark and Data Governance, no closing call-to-action in About, an inconsistent recent-posting cadence/topic strategy, and a business (not technical) undergraduate degree that could trip hard ATS degree-field filters at some employers. Recommended path forward: deepen into a Fabric-native AI engineering specialization while retaining Fabric/Power BI platform depth, rather than diluting or abandoning it.

SNAPSHOT

Profile Snapshot

- **Header:** present
- **About:** present
- **Top Skills:** 5 captured
- **Experience:** 4 captured
- **Education:** 2 captured
- **Skills Full List:** 28 captured
- **Certifications:** 3 captured
- **Featured Items:** 3 captured
- **Recent Posts:** 4 captured
- **Recommendations Received:** 3 captured
- **Recommendations Given:** 2 captured
- **Projects:** 6 captured
- **Missing:** none — all sections captured

FINDINGS

Section-by-Section Findings

About

- **CAREER COACH** — The about section and headline are framed entirely around Fabric/Power BI analytics engineering, but the two most recent projects (Aug 2026-Present, multi-agent Claude Code pipelines) and two skills_full_list entries (Claude Code, Multi-agent Systems) show active work in AI agent orchestration that is never mentioned in the headline or about. A reader has no way to tell if this is a hobby or a deliberate next direction.
- **CAREER COACH** — The about section closes on a certifications credential note (DP-700/DP-600) with no explicit call to action, no statement of what conversations or opportunities the person is open to.
- **CAREER COACH** — The opening line ("8+ years building the data infrastructure that powers strategic decisions") is serviceable but follows a common generic LinkedIn-about pattern; the strongest concrete proof points ($45M/$10M outcomes, both Fabric certifications) are buried in paragraphs three and four rather than leading.

Positioning & Narrative

- **CAREER COACH** — Headline, about section, top_skills, and current EPAM role all consistently point to the same target: Senior Analytics Engineer specializing in Microsoft Fabric, Power BI, Medallion architecture, semantic layer, and real-time analytics. Positioning is clear and unambiguous, not generic.

Experience

- **CAREER COACH** — All four roles are dense with quantified outcomes (35% latency reduction, 55% pipeline delivery cut, $45M opportunity generation, 38% churn reduction, 40% forecast accuracy gain). This is genuine impact framing, not task/responsibility listing, and is a real strength across the whole work history.
- **CAREER COACH** — All four experience entries follow an identical template (intro sentence, 4-5 metric-driven bullets, closing summary sentence). The uniformity across 8+ years and four different employers reads as formulaic rather than a distinct voice per role, which can look templated or AI-polished to a skeptical reader.
- **RECRUITER** — Every role includes specific, numeric outcomes (e.g., 35% latency reduction, 55% faster pipeline delivery, $31M and $210M divestiture support, $10M campaign revenue, 38% churn reduction, 40% forecast accuracy improvement). This is well above typical candidate quality for a skimming recruiter or hiring manager.
- **RECRUITER** — Each role description opens and closes with a narrative paragraph wrapped around a bulleted list, adding length that a recruiter skimming in 6 seconds is likely to skip past. The bullets themselves are strong; the framing prose adds scan friction.
- **RECRUITER** — The candidate's own experience section lists only 'Blackboard' for both Data Analyst roles (Nov 2019–Mar 2024). However, a received recommendation from Jason Bass states he worked with Felipe 'at Blackboard and then Anthology' -- a company name that never appears anywhere in the candidate's own profile sections. This is an unreconciled discrepancy a recruiter cross-referencing recommendations against experience could flag during verification.

Education

- **CAREER COACH** — A BBA (2012-2017) followed by a Cornell Data Analytics certificate (2019) provides a credible narrative bridge explaining the pivot from business administration into data analytics, reinforcing rather than undercutting the professional story.
- **RECRUITER** — Highest formal degree is a BBA in Business/Managerial Economics, not a technical/CS/data degree; the Cornell certificate program (2019) and Microsoft certifications are the technical credentialing. Some ATS configurations with hard degree-field filters for engineering roles could screen this out despite strong hands-on evidence elsewhere in the profile.

Skills & Keywords

- **CAREER COACH** — skills_full_list is otherwise a tight cluster of Fabric/analytics tools (Fabric, PySpark, Delta Lake, DAX, Azure Data Factory) but includes two outlier entries (Claude Code, Multi-agent Systems) tied only to the unaddressed side-project direction noted above, slightly diluting an otherwise focused skill signal.

- **RECRUITER** — top_skills (SQL, Python, Microsoft Fabric, Microsoft Power BI, Microsoft Azure) covers the broad platform terms but omits PySpark and Data Governance, both of which are emphasized repeatedly in about and experience and are terms a recruiter sourcing for a Fabric/data-engineering-heavy role would specifically search.
- **RECRUITER** — The 28-item list is mostly tight and relevant (PySpark, Medallion Architecture, Delta Lake, DAX, Data Modeling, ETL all reinforce the target role), but a few entries ('Claude Code', 'Multi-agent Systems', 'GitHub Actions') are tangential to the core Analytics/Data Engineer positioning and could dilute keyword focus for a recruiter scanning the skills section for Fabric/Power BI depth.

Certifications

- **CAREER COACH** — All three certifications (DP-700 issued Nov 2025, DP-600 issued Sep 2024, Power BI Data Analyst Associate issued Sep 2021) are current through 2027 and directly reinforce the Fabric/Power BI positioning stated in the headline and about section - strong credential-to-narrative alignment.
- **RECRUITER** — Three Microsoft certifications (Fabric Data Engineer Associate, Fabric Analytics Engineer Associate, Power BI Data Analyst Associate), all currently valid (none expired), and all directly relevant to the stated target role. This is a strong, verifiable credibility signal that a recruiter can validate quickly.

Featured

- **CAREER COACH** — Featured items curate the two most technically substantive GitHub projects plus the single best-performing post (199 reactions, ~44K impressions), showing good editorial instinct. However, the four visible recent posts swing from 405 to 47,685 impressions with no visible common thread, suggesting inconsistent posting cadence/topic strategy rather than a deliberate content plan.
- **RECRUITER** — Featured section includes two GitHub project write-ups plus a duplicate of the highest-performing recent post (199 reactions); all three are on-topic with the Fabric/data-engineering positioning, reinforcing consistency between headline, skills, and shown work.

Projects

- **RECRUITER** — Six personal projects strongly reinforce the Medallion/Fabric/PySpark keyword theme and demonstrate hands-on work beyond employer projects (NYC Taxi Medallion, Real-Time Stock Volatility, Market Basket Analysis mirroring the Blackboard work, Customer Churn). Two of the newest projects (Aug 2026) are about building this very Claude Code multi-agent audit pipeline, which is tangential to the core Fabric/Power BI/data-engineering positioning and could read as a scope departure to a recruiter narrowly sourcing for that specialty.

Recommendations

- **CAREER COACH** — Jason Bass's recommendation states Felipe "worked on my team for several years at Blackboard and then Anthology," but the experience section lists only "Blackboard" for both analyst roles (Nov 2019-Mar 2024) with no Anthology entry or reconciling note anywhere in the profile, leaving an unexplained gap between a recommender's account and the stated work history.
- **CAREER COACH** — Both recommendations given are specific and check out against the timeline: the note to the EPAM manager references DP-600 being earned, consistent with its Sep 2024 issue date and the Apr 2024-present EPAM tenure; the note to the Blackboard manager references 2020 revenue work, consistent with the Nov 2019-Jan 2022 Data Analyst role. No conflicts found, and reciprocated recommendations are themselves a positive relationship signal.
- **RECRUITER** — Three recommendations, all from named individuals with clear stated relationships, and all specific (naming actual projects, dashboards, and outcomes rather than generic praise). This is a strong, substantive credibility signal, not boilerplate.
- **RECRUITER** — Two thoughtful, specific recommendations given to former managers. This signals genuine professional engagement but has limited direct effect on searchability or skim assessment compared to recommendations received.

Posts & Visibility

- **RECRUITER** — Four posts spaced roughly weekly (1d, 1w, 2w, 3w old) show active, recent engagement rather than a stale profile. Topics (Fabric Eventstream vs. KQL trade-offs, MERGE/dedupe issues, memory-bound CSV loading, PySpark execution plans) directly reinforce the headline and top_skills positioning, which is exactly what a recruiter reviewing 'recent activity' wants to see.
- **RECRUITER** — Engagement is uneven but includes two standout posts (177 likes/47,685 impressions and 199 likes/44,089 impressions) well above typical organic reach, a strong credibility signal. However, comment counts are null or near-zero on 3 of 4 posts, suggesting the content generates views/likes but little discussion.

Header.Headline

- **RECRUITER** — Headline is specific, keyword-dense, and matches the current role and top skills ('Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics'). This is exactly the format a recruiter's keyword search rewards and a skim-reader can parse in under a second.

Header.Photo Present / Header.Banner Present

- **RECRUITER** — Both photo and banner are present, avoiding a common 6-second-skim red flag (missing photo is often read as an inactive or low-effort profile).

PRIORITIES

Prioritized Recommendations

High Impact

Reconcile the unaddressed AI/multi-agent pivot across About, skills, and projects (High Impact)

Career-coach flags this as the single biggest narrative fracture: the two most recent projects (Aug 2026-Present, multi-agent Claude Code pipelines) and two skills_full_list entries (Claude Code, Multi-agent Systems) show active work in AI agent orchestration that the headline and About never mention, leaving a reader unable to tell if it's a hobby or a deliberate next direction. Recruiter independently flags the same underlying signal from a screenability angle: those same skills and the two AI-tooling projects sit alongside the tightly-focused Fabric/Power BI keyword cluster, diluting keyword focus and potentially reading as a scope departure to a recruiter narrowly sourcing for a Fabric/data-engineering specialist.

Before

8+ years building the data infrastructure that powers strategic decisions. I specialize in designing scalable, AI-ready Medallion Lakehouses and real-time streaming pipelines on Microsoft Fabric, sitting at the intersection of data engineering and analytics. At EPAM Systems, I lead end-to-end data engineering initiatives: architecting Medallion architectures that process billions of records, building semantic models and Power BI solutions consumed by Finance, FP&A, Marketing, Strategy, and Sales, and pioneering a company-wide Data Governance framework that established 100% traceability across all financial reports. What sets my work apart is that I don't just build pipelines. I design the modeling layer, semantic layer, and data products that translate raw data into metrics stakeholders can trust and act on. That means dimensional modeling, clean semantic layers, and business-facing data products built to last. At Blackboard, that meant building the analytical foundation behind $45M in opportunity generation and $10M in campaign-driven revenue. I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack.

After

8+ years building the data infrastructure that powers strategic decisions. I specialize in designing scalable, AI-ready Medallion Lakehouses and real-time streaming pipelines on Microsoft Fabric, sitting at the intersection of data engineering and analytics. At EPAM Systems, I lead end-to-end data engineering initiatives: architecting Medallion architectures that process billions of records, building semantic models and Power BI solutions consumed by Finance, FP&A, Marketing, Strategy, and Sales, and pioneering a company-wide Data Governance framework that established 100% traceability across all financial reports. What sets my work apart is that I don't just build pipelines. I design the modeling layer, semantic layer, and data products that translate raw data into metrics stakeholders can trust and act on. That means dimensional modeling, clean semantic layers, and business-facing data products built to last. At Blackboard, that meant building the analytical foundation behind $45M in opportunity generation and $10M in campaign-driven revenue. I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack. Most recently, I've extended this data-product mindset into building multi-agent Claude Code pipelines that automate technical workflows end-to-end.

Resolve the Blackboard/Anthology company-name discrepancy (High Impact)

Career-coach flags an unexplained gap between Jason Bass's received recommendation, which states Felipe 'worked on my team for several years at Blackboard and then Anthology,' and the experience section, which lists only 'Blackboard' for both analyst roles (Nov 2019-Mar 2024) with no Anthology entry or reconciling note. Recruiter independently flags the identical discrepancy as a recruiter verification red flag: a company name that never appears in the candidate's own profile sections but does appear in a third-party recommendation is exactly the kind of inconsistency a recruiter cross-referencing recommendations against experience would catch during screening.

Before

Experience entries list company as "Blackboard" for both the Senior Data Analyst (Jan 2022 - Mar 2024) and Data Analyst (Nov 2019 - Jan 2022) roles, with no mention of Anthology anywhere in the profile.

After

Experience entries updated to note the entity change, e.g. company: "Blackboard (later Anthology)", so the profile and Jason Bass's external recommendation are consistent.

Vary experience-entry structure to reduce templated feel and skim friction (High Impact)

Career-coach flags that all four experience entries follow an identical template (intro sentence, 4-5 metric-driven bullets, closing summary sentence), which across 8+ years and four different employers reads as formulaic rather than a distinct voice per role. Recruiter independently flags the same narrative-wrapper pattern as scan friction for a 6-second skim: the bullets themselves are strong, but the framing prose around them adds length a skimming recruiter is likely to skip past, and recommends trimming intro/closing sentences so bullets are the first thing seen.

Before

Laid the foundation for my analytics career by modernizing reporting infrastructure at Renoboy S.A., transitioning the company from fragmented, manual Excel reports to a centralized, automated Power BI environment. - Delivered 8+ Power BI dashboards (Sales, Inventory, Cost & Profits, and Market Share) adopted across all senior management levels... - Replaced manual, Excel based sales forecasting (calculated using simple averages) with Python-built exponential smoothing and ARIMA models, standardizing the process across the sales team and improving forecast accuracy by 40%. - Migrated +8 Excel based reports across 4 business areas to automated Power BI dashboards, cutting manual reporting time by 60%. - Implemented a centralized SQL database consolidating 3 data sources with daily automated refresh, eliminating 70 hours of manual reconciliation. - Built 5 DAX-based KPI models used by 4 departments to track performance against targets. Transformed Renoboy's decision making culture from manual and reactive to automated and data-driven, establishing the data driven reporting infrastructure that shaped how leadership made decisions for years afterward.

After

Laid the foundation for my analytics career by modernizing reporting infrastructure at Renoboy S.A., transitioning the company from fragmented, manual Excel reports to a centralized, automated Power BI environment. - Delivered 8+ Power BI dashboards (Sales, Inventory, Cost & Profits, and Market Share) adopted across all senior management levels... - Replaced manual, Excel based sales forecasting (calculated using simple averages) with Python-built exponential smoothing and ARIMA models, standardizing the process across the sales team and improving forecast accuracy by 40%. - Migrated +8 Excel based reports across 4 business areas to automated Power BI dashboards, cutting manual reporting time by 60%. - Implemented a centralized SQL database consolidating 3 data sources with daily automated refresh, eliminating 70 hours of manual reconciliation. - Built 5 DAX-based KPI models used by 4 departments to track performance against targets.

Medium Impact

Add a closing call-to-action to the About section (Medium Impact)

Career-coach notes the About section closes on a certifications credential note (DP-700/DP-600) with no explicit call to action or statement of what conversations or opportunities Felipe is open to.

After

Always happy to connect with fellow data and analytics engineers or talk through Fabric architecture decisions - feel free to reach out.

Rebalance top_skills to surface higher-value searchable keywords (Medium Impact)

Recruiter notes top_skills (SQL, Python, Microsoft Fabric, Microsoft Power BI, Microsoft Azure) covers broad platform terms but omits PySpark and Data Governance, both emphasized repeatedly in About and experience and both terms a recruiter sourcing for a Fabric/data-engineering-heavy role would specifically search for.

Before

top_skills: ["SQL", "Python (Programming Language)", "Microsoft Fabric", "Microsoft Power BI", "Microsoft Azure"]

After

top_skills: ["SQL", "Python (Programming Language)", "Microsoft Fabric", "Microsoft Power BI", "PySpark"]

Develop a more deliberate recent-posts cadence and engagement strategy (Medium Impact)

Career-coach notes recent posts swing from 405 to 47,685 impressions with no visible common thread, suggesting inconsistent posting cadence/topic strategy rather than a deliberate content plan, even though featured items curate the best-performing content well. Recruiter separately notes comment counts are null or near-zero on 3 of 4 posts despite strong like/impression numbers, suggesting content generates views but little discussion, and suggests ending posts with a direct question to prompt comments.

Low Impact

Lead About with a concrete proof point instead of a generic opening line (Low Impact)

Career-coach notes the opening line ('8+ years building the data infrastructure that powers strategic decisions') is serviceable but follows a common generic LinkedIn-about pattern, while the strongest concrete proof points ($45M/$10M outcomes, both Fabric certifications) are buried in paragraphs three and four rather than leading.

Before

8+ years building the data infrastructure that powers strategic decisions.

After

I've architected Medallion Lakehouses processing billions of records and hold both Microsoft Fabric certifications (DP-700, DP-600) - here's how I turn raw data into numbers Finance and Sales actually trust.

BBA degree field may trigger hard ATS degree-filters for some engineering roles (Low Impact)

Recruiter notes the highest formal degree is a BBA in Business/Managerial Economics, not a technical/CS/data degree; the Cornell certificate program and Microsoft certifications are the technical credentialing. Some ATS configurations with hard degree-field filters for engineering roles could screen this out despite strong hands-on evidence elsewhere in the profile.

STRATEGY

Career Trajectory — 3-Year Outlook

Forward-looking strategic guidance on the candidate's career trajectory — distinct from the profile-polish recommendations above.

SWOT Analysis

Strengths

- Holds two current Fabric-specific certifications earned within the last 10 months (DP-700 Nov 2025, DP-600 Sep 2024) plus a 2021 Power BI cert, all valid through 2027.
- Clear seniority progression from Junior Data Analyst to Senior Analytics Engineer in under 8 years, including leading a team of 3 at EPAM.
- Consistent project cadence since Jan 2025 covering Medallion architecture, streaming, and ML - a deliberate skill-building pattern, not one-off tinkering.
- Quantified, name-brand business impact on record: $45M opportunity generation, support for $210M and $31M divestiture decisions at Blackboard.

Opportunities

- AI engineer job postings grew roughly 109% year-over-year with average pay near $206K and a 3.2:1 demand-to-supply gap, a plausible on-ramp given his Fabric, Python, and Spark base plus new agent-project cadence.
- AI governance demand is up about 150% year-over-year per LinkedIn's 2026 Skills on the Rise data, directly adjacent to the governance framework he already built at EPAM.
- Enterprises already on Microsoft 365/Power BI see meaningfully lower total cost consolidating onto Fabric, keeping his specialization in steady demand at a multi-client consultancy like EPAM.

Weaknesses

- No Databricks, dbt, Snowflake, or non-Microsoft cloud exposure anywhere in skills_full_list or projects - entirely single-vendor platform depth.
- No containerization or orchestration skills (Docker, Kubernetes) despite these increasingly appearing in adjacent AI/MLOps job requirements.
- The new multi-agent/Claude Code project work is a month old, isolated to personal GitHub repos, with no certification or employer-attributed production experience behind it yet.
- Governance experience is data-catalog and lineage focused, not the emerging AI-governance specialization (model risk, compliance frameworks) that market data shows growing fastest.

Threats

- Databricks remains the market leader for heavy ML and AI-engineering workloads, so Fabric-only depth without any Databricks exposure narrows him out of Databricks-standardized employers.
- Platform market-share data is genuinely mixed across sources - one source shows Fabric at roughly 12% share versus Databricks 5%, while another calls Fabric the fastest-growing analytics platform ever - a real signal but not one to over-index on either direction.
- Production agentic-AI and MLOps roles increasingly expect Kubernetes/Docker (present in over a quarter of LangChain-ecosystem job postings), an area where he currently has zero footprint.

Recommended Paths

Stay IC, deepen into a Fabric-native AI engineering specialization

Rationale:

Builds directly on his DP-700/DP-600 base, existing governance framework, and new agent-project cadence, positioning him inside the fast-growing AI engineer demand curve without abandoning his platform depth.

Tradeoffs:

Fabric-anchored AI tooling is a narrower market than open-source or Databricks-based AI engineering, and he would need to prove production-grade (not personal-project) agentic work to be credible for a real pivot.

Add Databricks depth to become a platform-agnostic analytics/data engineer

Rationale:

Databricks still leads heavy ML workloads, and diversifying beyond a single vendor widens the employer pool beyond Microsoft-first shops - especially relevant inside a multi-client consultancy like EPAM.

Tradeoffs:

Splitting focus across two ecosystems costs real time and dilutes the tightly-differentiated 'Fabric specialist' positioning that currently makes this profile stand out.

Pursue a formal analytics/data leadership track

Rationale:

He has already led a team of 3 and mentored juniors across three separate employers, a repeated pattern, and Analytics Engineer comp is strong even before a management premium.

Tradeoffs:

A people-management track would slow the hands-on technical depth-building (Fabric certs, agent projects) that currently differentiates him, and the profile shows no formal leadership title or credential signaling readiness yet.

Skill & Certification Recommendations

Databricks exposure (project work or Databricks Certified Data Engineer Associate) (Medium Priority)

Broadens beyond single-vendor Fabric dependency, and Databricks still leads heavy ML/AI workloads in the market.

AI governance / responsible AI framework knowledge (High Priority)

Extends his existing EPAM governance-framework work into the fastest-growing specialization tracked in 2026 market data rather than starting a new skill from zero.

Production-grade deployment of the multi-agent work, plus basic Docker/Kubernetes literacy (High Priority)

Agent-framework jobs are increasingly production infrastructure roles, and his current agent work is unproven beyond personal repos.

LangChain specifically (Low Priority)

He is already building multi-agent systems directly with Claude Code; the infra and governance skills matter more than adopting this particular library, so it is low urgency for him specifically.

dbt or an equivalent transformation tool outside Power BI/DAX (Low Priority)

He already owns semantic-layer depth inside Fabric; this only becomes worthwhile if he commits to the Databricks/platform-agnostic path rather than as a universal priority.

OODA Execution Loop

Observe

Each quarter, scan Analytics Engineer, AI Engineer, and Data Engineer postings for how often Fabric is paired with Databricks or Kubernetes requirements, check LinkedIn's Skills on the Rise report for AI governance movement, and track which of his own posts get the most engagement.

→

Orient

Compare what surfaced against the skill gaps named in the SWOT weaknesses above (Databricks, Kubernetes, AI governance) and re-check the SWOT findings to see which weaknesses are closing and which threats (Fabric vs. Databricks share, production-infra requirements) are intensifying.

→

Decide

Pick which of the three recommended paths to lean into that quarter based on which market signal moved most since the last cycle, rather than chasing all three in parallel indefinitely.

→

Act

Ship one concrete artifact per quarter proving the chosen direction - a Databricks project, a production-deployed multi-agent pipeline with monitoring, or an expanded leadership scope at EPAM - and only update the headline/about section once that evidence is real, so the profile doesn't outrun the work again.

ROADMAP

Next Steps

This Week

- Add the one-sentence AI/multi-agent pivot acknowledgment to the About section (Reconcile the unaddressed AI/multi-agent pivot across About, skills, and projects)
- Update the two Blackboard experience entries to note 'Blackboard (later Anthology)' (Resolve the Blackboard/Anthology company-name discrepancy)
- Trim or drop the closing summary sentence on the Renoboy and early Blackboard experience entries so bullets stand on their own (Vary experience-entry structure to reduce templated feel and skim friction)
- Add a closing call-to-action line to the About section (Add a closing call-to-action to the About section)
- Swap 'Microsoft Azure' for 'PySpark' in the top_skills pinned list (Rebalance top_skills to surface higher-value searchable keywords)

- Rewrite the About section's opening line to lead with a concrete proof point (Lead About with a concrete proof point instead of a generic opening line)

This Month

- Decide whether skills_full_list and project ordering should deprioritize the AI-tooling entries or lean further into a dual Fabric-plus-AI-agent positioning, once the About pivot sentence from short-term is live (Reconcile the unaddressed AI/multi-agent pivot across About, skills, and projects)
- Draft a deliberate posting topic/cadence plan and experiment with closing questions on new posts to increase comment engagement (Develop a more deliberate recent-posts cadence and engagement strategy)

This Quarter

- Over the quarter, publish posts against the new cadence/topic plan and monitor whether comment engagement improves to confirm the strategy is working (Develop a more deliberate recent-posts cadence and engagement strategy)
- Monitor whether the BBA degree field is triggering ATS screen-outs at target companies when applying, since no direct profile edit resolves this gap (BBA degree field may trigger hard ATS degree-filters for some engineering roles)