LinkedIn Profile Audit Report

Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics

Generated August 30, 2026

OVERVIEW

Executive Summary

This is a strong, well-positioned profile: the headline, about section, top_skills, certifications (DP-700, DP-600, Power BI Data Analyst), and four roles of consistently quantified achievements (35% latency reduction, $45M in opportunity generation, $10M in campaign-driven revenue, billions of records processed) all reinforce a coherent 'Microsoft Fabric analytics engineer' narrative that would survive both a recruiter's 6-second skim and an ATS keyword search. Substantive, specific recommendations received, active and well-engaged recent posts (up to 199 likes / 44,085 impressions), and a deliberate portfolio of six self-directed GitHub projects further back up the story with genuine credibility. The two headline issues worth fixing are an unreconciled Blackboard/Anthology company-name discrepancy that a recruiter cross-checking the timeline against a recommendation could flag, and a skills/branding tension around the profile's real, recent AI and multi-agent work: it's currently invisible in the headline/about/top_skills while simultaneously sitting at the top of skills_full_list ahead of the core Fabric/Power BI/SQL keywords that should lead. Secondary gaps include a missing call-to-action at the end of the about section and no recommendation covering the current, most senior EPAM role. Recommended path forward: deepen into applied AI/agentic engineering on top of the existing Fabric base, while framing it as an extension of the current specialization rather than a pivot away from it.

SNAPSHOT

Profile Snapshot

- **Header:** present
- **About:** present
- **Top Skills:** 5 captured
- **Experience:** 4 captured
- **Education:** 2 captured
- **Skills Full List:** 28 captured
- **Certifications:** 3 captured
- **Featured Items:** 3 captured
- **Recent Posts:** 4 captured
- **Recommendations Received:** 3 captured
- **Recommendations Given:** 2 captured
- **Projects:** 6 captured
- **Missing:** none — all sections captured

FINDINGS

Section-by-Section Findings

About

- **CAREER COACH** — The about section has a strong, specific hook ('8+ years building the data infrastructure that powers strategic decisions') and concrete proof points, but it ends abruptly on the certification list with no call to action inviting connection, conversation, or opportunity discussion.

Positioning & Narrative

- **CAREER COACH** — The headline reads 'Senior Analytics Engineer' while the current role in the experience section is titled 'Senior Data Analytics Engineer.' The overall Fabric/Power BI/Medallion thread is otherwise consistent across headline, about, and experience, so this is a minor wording drift rather than a real contradiction.
- **CAREER COACH** — Jason Bass's recommendation states he worked with Felipe 'at Blackboard and then Anthology,' but 'Anthology' never appears anywhere in the experience section, which lists only 'Blackboard' for both roles spanning Nov 2019-Mar 2024. Blackboard and Anthology merged in 2021, so this is plausibly a real rebrand, but the profile never reconciles the naming, leaving a discrepancy a recruiter cross-checking employer history could flag.
- **CAREER COACH** — Positioning is clear and consistent: the profile targets senior/staff-level Microsoft Fabric-focused analytics engineering roles, reinforced by matching headline keywords, about-section specifics, and certifications (DP-700, DP-600). This is focused rather than generic, which is a genuine strength.

Experience

- **CAREER COACH** — All four roles list outcome-based bullets with concrete metrics rather than pure task lists (e.g., 35% latency reduction, 100+ tables served, $45M in opportunity generation, 80% cross-sell accuracy). This achievement/impact framing is one of the strongest parts of the profile.
- **CAREER COACH** — The Junior Data Analyst (Renoboy S.A.) bullets lean more toward task/responsibility language ('Delivered 8+ dashboards,' 'Replaced manual Excel reporting') than the outcome framing used in later roles. Understandable for an entry-level role, but the claim that dashboards were 'adopted across all senior management levels' currently has no quantified business result attached.
- **RECRUITER** — Achievements across all four roles are consistently quantified (35% latency reduction, 40% faster response, $45M in opportunity generation, $10M in campaign-driven revenue, 100+ tables, 20+ reports, 80% accuracy, billions of records). This is well above the norm and gives a skimming recruiter concrete signal instead of vague responsibility statements.
- **RECRUITER** — The EPAM Systems and Senior Data Analyst (Blackboard) entries each run 6 bullet points plus opening/closing prose paragraphs. It's bulleted rather than a true wall of text, but it's dense enough that a recruiter skimming for ~6 seconds may only register the first 1-2 bullets before moving on.
- **RECRUITER** — No unexplained employment gaps: Renoboy S.A. (Dec 2017-Nov 2019) -> Blackboard Data Analyst (Nov 2019-Jan 2022) -> Blackboard Senior Data Analyst (Jan 2022-Mar 2024) -> EPAM Systems (Apr 2024-Present) form a continuous timeline. Job titles are specific and show a clear, legible progression (Junior -> Data Analyst -> Senior Data Analyst -> Senior Data Analytics Engineer).

Education

- **CAREER COACH** — A Business/Managerial Economics BBA plus a Cornell Data Analytics certificate (2019) gives credible grounding to the about section's claims of translating data into business-trusted metrics; the business degree is a genuine differentiator versus a purely technical background.
- **RECRUITER** — BBA in Business/Managerial Economics plus a Cornell Data Analytics certificate provide reasonable context for an analytics career path; no red flags (no unexplained gaps, dates are present and consistent with the Dec 2017 start of the first analyst role).

Skills & Keywords

- **CAREER COACH** — skills_full_list and projects show real, recent (2026) hands-on work in AI and multi-agent systems ('Artificial Intelligence (AI)', 'Multi-agent Systems', 'Claude Code' skills, plus two dedicated 2026 GitHub projects), but none of this appears in the headline, about section, or top_skills. A differentiated, timely skill area is currently invisible to anyone skimming the profile.
- **RECRUITER** — The 5 pinned top_skills (SQL, Python, Microsoft Fabric, Microsoft Power BI, Microsoft Azure) are all core, current, and consistent with the headline and about section — strong searchability signal with no dilution.

- **RECRUITER** — The 28-item skills list opens with Artificial Intelligence (AI), GitHub Actions, Claude Code, and Multi-agent Systems, while the core ATS-searchable keywords that match the headline and top_skills (SQL, Python, Microsoft Azure, Microsoft Fabric, Microsoft Power BI) sit at the very bottom (positions 24-28). If LinkedIn surfaces skills in this order before 'show all,' a recruiter's first glance would see AI-tooling terms rather than the Fabric/Power BI/SQL signal the rest of the profile is built around.

Certifications

- **CAREER COACH** — All three certifications are current, directly relevant, and non-expired (soonest expiry Sep 2027), and they directly back up the headline's Fabric positioning (DP-700, DP-600) plus the earlier Power BI Data Analyst credential.
- **RECRUITER** — Three Microsoft certifications directly relevant to the stated role (Fabric Data Engineer Associate, Fabric Analytics Engineer Associate, Power BI Data Analyst Associate), all current (not expired), and explicitly cross-referenced in the about section. This is a strong credibility signal for both ATS keyword matching and human review.

Featured

- **CAREER COACH** — Recent posts show real, growing engagement on technical Fabric/PySpark topics (e.g., 199 likes / 44,085 impressions on an execution-plan post; 177 likes / 47,682 impressions on a memory-management post). This is genuine thought-leadership traction that reinforces the profile's stated positioning.

Projects

- **RECRUITER** — Six projects are described, all with clear tech stacks and GitHub links. The four Fabric/Power BI/Python projects (NYC Taxi Medallion, Real-Time Stock Volatility, Customer Churn Prediction, Market Basket Analysis) tightly reinforce the stated target role. The two most recent projects (LinkedIn Profile Audit multi-agent pipeline, Multi-Agent Blog Digest, both Aug 2026-Present) are AI-agent/automation builds that are adjacent but not core to the 'Analytics Engineer / Fabric / Power BI' positioning — useful as an AI-fluency signal but could slightly dilute a recruiter's read of the primary specialization if they're keyword-scanning strictly for Fabric/Power BI project work.

Recommendations

- **CAREER COACH** — All three recommendations received are from the Blackboard era (2019-2024). There is no recommendation received covering the current EPAM role (Apr 2024-present), despite EPAM being the most senior, most recent position and the one the headline and about section foreground most.
- **CAREER COACH** — Two thoughtful recommendations given (to an EPAM manager and a former Blackboard manager) show active, reciprocal engagement with the platform. This has limited direct bearing on Felipe's own positioning but supports an overall impression of an engaged, relationship-oriented professional, consistent with the recommendations received.
- **RECRUITER** — Jason Bass's recommendation states Felipe 'worked on my team for several years at Blackboard and then Anthology,' but the experience section only lists 'Blackboard' for both Data Analyst roles (Nov 2019 - Mar 2024) — Anthology never appears as a distinct entry or note. A recruiter cross-referencing the recommendation against the experience timeline would see an unreconciled company name.
- **RECRUITER** — All three recommendations are specific and substantive (named projects like the GTM Power BI report and Market Basket Analysis Dashboard, named companies, specific outcomes) rather than generic boilerplate — a genuine credibility signal, not just a headcount of recommendations.
- **RECRUITER** — Two recommendations given (to a manager at EPAM and a manager at Blackboard) show reciprocal engagement on the platform. Not a primary screening factor, but a minor positive signal of an active, networked profile.

Posts & Visibility

- **RECRUITER** — 4 recent posts, well spaced and recent (19h, 1w, 2w, 3w), signaling active, ongoing engagement rather than a stale profile. All four are technical deep-dives on Fabric/PySpark/streaming topics, directly reinforcing the headline and top_skills positioning rather than drifting off-topic.
- **RECRUITER** — Engagement is strong on two posts (177 likes/47,682 impressions; 199 likes/44,085 impressions) but comment counts are null/near-zero on all four posts (null, null, 2, null) despite the high-impression posts. Likes and impressions are a credible visibility signal, but the near-absence of comments suggests the content isn't sparking discussion, which weakens the reach/authority signal somewhat compared to a post with active comment threads.

Header.Headline

- **RECRUITER** — Headline is specific and keyword-dense ('Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics'), directly matching the terms a recruiter sourcing for this role would search. Photo and banner are both present, removing two common 6-second-skim red flags.

Missing Fields

- **RECRUITER** — missing_fields is empty and every schema section (header, about, top_skills, experience, education, skills_full_list, certifications, featured_items, recent_posts, recommendations_received, recommendations_given, projects) is populated with real data — nothing had to be excluded from this review for lack of data.

PRIORITIES

Prioritized Recommendations

High Impact

Reconcile the Blackboard / Anthology company-name discrepancy (High Impact)

Both agents independently flag the same issue from two angles: career-coach notes that Jason Bass's recommendation states he worked with Felipe 'at Blackboard and then Anthology,' but 'Anthology' never appears in the experience section, which lists only 'Blackboard' for both roles spanning Nov 2019-Mar 2024. Recruiter flags this as a cross-reference red flag a detail-oriented recruiter could catch when checking the recommendation against the employment timeline. Blackboard and Anthology merged in 2021, so this is very likely a real rebrand the profile simply never reconciles - but as written it reads as an unexplained inconsistency.

Before

Experience entries list the company simply as "Blackboard" for both the Senior Data Analyst (Jan 2022 - Mar 2024) and Data Analyst (Nov 2019 - Jan 2022) roles, with no reference to Anthology anywhere in the experience section.

After

Add a brief parenthetical to both Blackboard experience entries noting the 2021 rebrand/merger, e.g. "Blackboard (part of Anthology from 2021)."

Rebalance the AI / multi-agent signal without diluting the core Fabric/Power BI keyword positioning (High Impact)

career-coach flags that real, recent (2026) AI and multi-agent systems work (skills_full_list entries, plus two dedicated 2026 GitHub projects) is invisible in the headline, about section, and top_skills, despite being a timely, differentiated skill area. Recruiter flags the inverse side of the same issue from a searchability angle: skills_full_list currently opens with AI/automation terms (Artificial Intelligence (AI), GitHub Actions, Claude Code, Multi-agent Systems) while the core ATS-searchable keywords that match the headline and top_skills (SQL, Python, Microsoft Fabric, Microsoft Power BI, Microsoft Azure) sit buried at positions 24-28; recruiter separately notes the two AI-agent projects, while genuine differentiators, could slightly dilute a recruiter's read of the primary specialization if left unframed. Both agents are pointing at the same underlying tension in the skills/projects sections: the AI work is worth surfacing deliberately, but not at the expense of the core keyword signal the rest of the profile is built around.

Before

skills_full_list currently opens with "Artificial Intelligence (AI), GitHub Actions, Claude Code, Multi-agent Systems, Microsoft Fabric Notebooks..." with the core keywords (SQL, Microsoft Fabric, Microsoft Power BI, Microsoft Azure, Python (Programming Language)) at positions 24-28. The about section currently ends: "...covering both the engineering and analytics layers of the Fabric stack." with no mention of the AI/multi-agent project work.

After

Reorder skills_full_list to lead with the core keywords ("SQL, Microsoft Fabric, Microsoft Power BI, Microsoft Azure, Python (Programming Language), ..."), moving AI/automation terms (Artificial Intelligence (AI), Claude Code, Multi-agent Systems, GitHub Actions) further down as supplementary skills. Add "Multi-agent AI systems" to top_skills. Extend the about section's closing with one sentence, e.g.: "I'm also applying that same platform-building rigor to hands-on multi-agent AI systems work (the LinkedIn Profile Audit and Multi-Agent Blog Digest Claude Code pipelines), extending my data platform expertise into applied AI tooling."

Medium Impact

Add a closing call to action to the About section (Medium Impact)

career-coach notes the about section has a strong, specific hook ("8+ years building the data infrastructure that powers strategic decisions") and concrete proof points, but ends abruptly on the certification list with no call to action inviting connection, conversation, or opportunity discussion.

Before

About section currently ends: "I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack."

After

Add a closing line: "If you're building AI-ready data platforms on Fabric or thinking through data governance at scale, I'd welcome the conversation - feel free to connect."

Get a recommendation covering the current EPAM role (Medium Impact)

career-coach notes all three recommendations received are from the Blackboard era (2019-2024). There is no recommendation covering the current, most senior EPAM role (Apr 2024-present), despite EPAM being the position the headline and about section foreground most.

Low Impact

Trim dense experience bullets for a faster skim (Low Impact)

recruiter notes the EPAM Systems and Senior Data Analyst (Blackboard) entries each run 6 bullet points plus opening/closing prose; dense enough that a recruiter skimming for ~6 seconds may only register the first 1-2 bullets before moving on.

Align headline title wording with the current experience title (Low Impact)

career-coach notes the headline reads "Senior Analytics Engineer" while the current role in the experience section is titled "Senior Data Analytics Engineer" - a minor wording drift rather than a real contradiction, since the surrounding Fabric/Power BI/Medallion positioning is otherwise consistent.

Add a quantified outcome to the Junior Data Analyst (Renoboy) bullets (Low Impact)

career-coach notes the Renoboy S.A. bullets lean toward task/responsibility language rather than the outcome framing used in later roles, and the claim that dashboards were "adopted across all senior management levels" currently has no quantified business result attached.

End posts with a question to prompt comments (Low Impact)

recruiter notes recent posts show strong likes/impressions (up to 199 likes / 44,085 impressions) but near-zero comments across all four posts, which weakens the reach/authority signal compared to posts that spark active discussion.

STRATEGY

Career Trajectory — 3-Year Outlook

Forward-looking strategic guidance on the candidate's career trajectory — distinct from the profile-polish recommendations above.

SWOT Analysis

Strengths

- Holds DP-700 and DP-600 certifications, covering both the Fabric engineering and analytics layers, plus a 2021 Power BI Data Analyst credential - full-stack Fabric coverage.
- Experience bullets consistently quantify outcomes: 35% latency reduction, 40% faster incident response, $45M in opportunity generation, $10M in campaign revenue.
- Six self-directed GitHub projects span Medallion architecture, real-time streaming, and ML, showing deliberate skill-building beyond day-job scope, not one-off tutorials.
- Currently leads a 3-person data engineering team at EPAM while remaining a hands-on architect, giving early management exposure without abandoning technical depth.

Opportunities

- Microsoft Fabric grew to 31,000+ customers and $2B ARR, up 60% year over year, and runs 20-40% cheaper than Databricks for comparable workloads.
- AI governance demand grew 150% year over year with a 25-35% salary premium; his existing EPAM data-governance framework is a direct stepping stone.
- Analytics Engineer median base salary ($158K) now edges out Data Engineer ($153K), rewarding his semantic-layer and dimensional-modeling specialization over generic pipeline work.
- Fabric and Databricks now share data bidirectionally without copying (GA June 2026), creating demand for engineers fluent in both platforms rather than one.

Weaknesses

- No Databricks, Snowflake, or non-Microsoft lakehouse platform appears anywhere in the skills list or project portfolio, leaving lakehouse tooling entirely single-vendor.
- No LLM or agentic-AI tooling (LangChain, RAG, vector databases) appears in the skills list, despite two 2026 projects building multi-agent Claude Code pipelines.
- All three certifications and nearly every listed skill are Microsoft-specific; no cloud-agnostic or multi-cloud credential signals platform flexibility to non-Microsoft-stack employers.
- Leadership scope tops out at a 3-person team, with no evidence yet of budget ownership, hiring, or cross-team org design needed for a manager-track role.

Threats

- Databricks still leads Fabric on advanced MLOps (MLflow, Feature Store, Model Serving), an area with zero current profile presence if enterprise ML work shifts platform.
- Recurring reports, dashboard updates, and KPI summaries face rising automation pressure; junior-analyst-style reporting tasks show 55-70% AI exposure risk industry-wide.
- One industry account describes senior analysts who took manager roles losing roughly 18 months before reverting to IC work - a real risk for any premature management pivot.

Recommended Paths

Deepen into applied AI/agentic engineering on top of the existing Fabric base.

Rationale:

He already has two 2026 personal projects building multi-agent Claude Code pipelines and the EPAM role description explicitly points at 'future Generative AI and LLM integrations'; agentic AI skills currently carry the steepest salary premiums in the market.

Tradeoffs:

Requires learning an entirely new tooling stack (vector databases, LangGraph, RAG) largely outside current day-job scope, and risks diluting the clean 'Fabric specialist' narrative if not framed as an extension rather than a pivot.

Pursue an Analytics Manager / team-lead track.

Rationale:

He is already leading a 3-person team at EPAM, was 'promoted to lead' initiatives at Blackboard, and multiple recommendations praise his mentoring; analytics/data roles including management layers are projected to grow strongly this decade.

Tradeoffs:

Search evidence shows a real risk of stalling roughly 18 months and reverting to IC work if the transition happens before management scope is deliberately expanded and tested first.

Add Databricks depth alongside Fabric to stay platform-agnostic.

Rationale:

Enterprises increasingly run Fabric and Databricks together with GA bidirectional data sharing as of June 2026, and Databricks retains a real MLOps edge Fabric currently lacks, widening his addressable market beyond Microsoft-only shops.

Tradeoffs:

Splitting depth across two lakehouse platforms takes time away from going deeper on Fabric-native differentiation, which is what his current headline and DP-700/DP-600 combo are built around.

Skill & Certification Recommendations

LangChain/LangGraph and RAG fundamentals (Medium Priority)

Builds directly on his existing Python and multi-agent Claude Code project work and captures the steepest AI salary premiums currently in the market.

AI governance practices (e.g., familiarity with the NIST AI RMF) (High Priority)

Directly extends the data-governance framework he already built at EPAM into the fastest-growing governance niche, with minimal new tooling to learn.

Databricks Certified Data Engineer Associate (Medium Priority)

Closes his single-vendor platform gap and positions him for the growing share of enterprises running Fabric and Databricks together.

Formal people-management training (Low Priority)

Only worth pursuing if he actively chooses the Analytics Manager path; premature without first testing management on a larger scoped team.

Generic AI/prompt-engineering certificates unrelated to a specific LLM-ops direction (Low Priority)

Not worth his time - his hands-on multi-agent project portfolio already demonstrates more than most generic certificates would signal.

OODA Execution Loop

Observe

Each quarter, review EPAM project scope changes, Fabric/Databricks feature-parity announcements (e.g., OneLake-Databricks bidirectional sharing), which LinkedIn posts land best (his PySpark/memory-management posts already pull 40K+ impressions), and whether AI-governance or AI-engineering work is opening up internally.

→

Orient

Compare what was observed against this SWOT: is the Databricks/MLOps weakness closing or widening, is AI governance actually creating internal opportunities at EPAM, and is his 3-person team growing (a signal favoring the manager path) or staying flat (a signal favoring staying IC)?

→

Decide

At each quarterly checkpoint, choose one lever to invest in rather than three - e.g., ship one more portfolio project (Databricks-based or LangChain-based), pursue one governance-adjacent initiative at EPAM, or explicitly request expanded management scope.

→

Act

Execute that single chosen action within the quarter, then update the LinkedIn about section, top_skills, and projects to reflect the new evidence only after it exists (never in advance of it), and re-check the SWOT findings at the next quarterly review. Grounded in profile data (skills_full_list, certifications, experience, projects) and 6 web searches on market demand, run this session.

ROADMAP

Next Steps

This Week

- Add a parenthetical noting the 2021 Blackboard/Anthology rebrand to both Blackboard experience entries (Reconcile the Blackboard / Anthology company-name discrepancy).
- Add a closing call-to-action line to the About section (Add a closing call to action to the About section).

- Reorder skills_full_list so core keywords (SQL, Microsoft Fabric, Microsoft Power BI, Microsoft Azure, Python) lead and AI/automation terms move further down (Rebalance the AI / multi-agent signal).
- Update the headline's title fragment to match the current experience title, "Senior Data Analytics Engineer" (Align headline title wording with the current experience title).

This Month

- Add "Multi-agent AI systems" to top_skills and one sentence to the About section referencing the applied AI project work (Rebalance the AI / multi-agent signal).
- Trim the EPAM and Senior Data Analyst (Blackboard) bullet lists to the top 3-4 highest-impact, most quantified bullets each (Trim dense experience bullets for a faster skim).
- Add a quantified outcome (e.g., time saved per reporting cycle, or a specific decision enabled) to the Renoboy S.A. bullets (Add a quantified outcome to the Junior Data Analyst (Renoboy) bullets).
- End future LinkedIn posts with a direct question to prompt comments rather than likes alone (End posts with a question to prompt comments).

This Quarter

- Request a recommendation from an EPAM colleague or manager covering the current, most senior role (Apr 2024-present) (Get a recommendation covering the current EPAM role).