LinkedIn Profile Audit Report

Senior Data Engineer | Microsoft Fabric · PySpark · Python · SQL | Medallion Lakehouse · Real-Time Streaming | AI & RAG Data Systems

Generated September 21, 2026

OVERVIEW

Executive Summary

This is a strong, coherent profile: the headline, About section, and every experience bullet reinforce one consistent Fabric-centric Medallion Lakehouse story, achievements are quantified almost everywhere (35% latency reduction, 55% defect-rate cut, $45M and $10M revenue figures), current Microsoft certifications back the stated specialization, photo and banner are present, there are no employment gaps, and posting cadence is active and topically on-brand. The headline issues are structural rather than narrative: the pinned top_skills list and Featured items lag behind the profile's newest AI/RAG and Databricks-flavored work, all three received recommendations are frozen at the Blackboard era with none validating the current, most senior EPAM role, a company-naming mismatch (Blackboard vs. Anthology) sits unreconciled between the experience section and one recommendation, and a non-technical bachelor's degree carries some ATS risk that the certifications and project links already largely offset. None of these are narrative problems - they are gaps between what the profile says and what it visibly shows. Recommended path: stay an individual contributor and deepen into AI/Agentic Data Engineering, building on the existing Fabric and Databricks lakehouse depth, while closing the recommendation and keyword-alignment gaps above.

SNAPSHOT

Profile Snapshot

- **Header:** present
- **About:** present
- **Top Skills:** 5 captured
- **Experience:** 4 captured
- **Education:** 2 captured
- **Skills Full List:** 34 captured
- **Certifications:** 3 captured
- **Featured Items:** 2 captured
- **Recent Posts:** 7 captured
- **Recommendations Received:** 3 captured
- **Recommendations Given:** 2 captured
- **Projects:** 6 captured
- **Missing:** none — all sections captured

FINDINGS

Section-by-Section Findings

About

- **CAREER COACH** — The opening line ('8+ years building the data infrastructure that powers strategic decisions') is serviceable but leans on a somewhat generic phrase rather than a concrete number up front. More notably, the about section ends on certification facts with no call to action inviting connection, conversation, or a look at the featured work.

- **CAREER COACH** — Beyond the hook, the body is well above average for specificity: named platforms (Fabric, Power BI), named business functions (Finance, FP&A, Marketing, Strategy, Sales), and named dollar figures ($45M, $10M) replace generic filler language throughout.

Positioning & Narrative

- **CAREER COACH** — Headline, about section, and experience entries tell one coherent story: Fabric-centric Medallion Lakehouse engineering with consistently carried quantified outcomes across both EPAM and Blackboard roles. This is a genuine strength.
- **CAREER COACH** — Jason Bass's recommendation states Felipe worked 'at Blackboard and then Anthology,' but the experience section lists only 'Blackboard' for both analyst roles (Nov 2019-Mar 2024), with no Anthology entry or rebrand note anywhere in the profile to reconcile the two names.
- **CAREER COACH** — Headline and about section position Felipe exclusively as a Microsoft Fabric specialist, yet the two most recent flagship projects (Aug-Sep 2026) are built on Databricks rather than Fabric, creating tension between the stated platform specialization and the newest, most visible project work.

Experience

- **CAREER COACH** — All four roles use consistent outcome-based bullet framing with quantified metrics (35% latency reduction, 55% defect-rate cut, 40% forecast-accuracy improvement, $45M and $10M revenue figures, hours saved) rather than pure task/responsibility listing. This is one of the strongest parts of the profile.
- **RECRUITER** — Nearly every bullet across all four roles is quantified with specific metrics or dollar figures (e.g., 35% latency reduction, 55% defect cut, 38% churn reduction, $45M opportunity generation, $210M sale), which is a strong credibility signal for a skimming recruiter.
- **RECRUITER** — No unexplained gaps: roles run continuously from Dec 2017 to present across Renoboy S.A, Blackboard (two titles), and EPAM Systems.
- **RECRUITER** — Each role uses a short intro sentence plus scannable bullets rather than dense paragraphs, and job titles (Senior Data Engineer, Senior Data Analyst, Data Analyst, Junior Data Analyst) are specific and show clear progression, both of which suit a 6-second skim.

Education

- **CAREER COACH** — A Business Administration degree plus a Cornell data analytics certificate underpins the 'business acumen + technical depth' narrative that both the about section and multiple recommendations (Claudio Warneboldt, Jason Bass) independently echo - a consistent, mutually reinforcing signal.
- **RECRUITER** — The bachelor's degree is in Business Administration (Business/Managerial Economics) rather than a technical field, plus a 2019 Cornell Data Analytics certificate. A rigid ATS filter screening for a technical degree could deprioritize the profile despite the strong offsetting hands-on experience and DP-700/DP-600 certifications.

Skills & Keywords

- **CAREER COACH** — The headline leads with 'AI & RAG Data Systems' as a core positioning pillar, but the pinned top_skills list (SQL, Python, PySpark, Microsoft Fabric, Data Engineering) contains nothing related to AI, RAG, or multi-agent systems.
- **RECRUITER** — Headline and top_skills strongly echo each other on the core searchable terms recruiters would use (Microsoft Fabric, PySpark, Python, SQL, Data Engineering), and the about section reinforces them with Medallion Lakehouse, real-time streaming, and AI/RAG language.
- **RECRUITER** — Databricks appears heavily in the about section, in 2 of 6 projects, and in several recent_posts, but is absent from both the headline and the 5-slot top_skills list. A recruiter sourcing specifically on 'Databricks' as a standalone search term may not surface this profile.
- **RECRUITER** — The 34-item skills list is current and relevant throughout (engineering, governance, BI, and AI/agent tooling); there is no outdated or filler technology diluting the signal.

Certifications

- **CAREER COACH** — All three certifications are Microsoft-only (DP-700, DP-600, PL-300), and the about section explicitly names the two Fabric certs, reinforcing a tight, consistent single-platform specialization narrative. There is no certification backing the Databricks/Apache Spark skills also listed in skills_full_list.
- **RECRUITER** — Three Microsoft certifications (Fabric Data Engineer Associate, Fabric Analytics Engineer Associate, Power BI Data Analyst Associate) are all current (not expired) and directly relevant to the stated positioning, giving strong ATS/keyword credibility.

Featured

- **CAREER COACH** — The two pinned featured items are both Fabric-only medallion/streaming projects. The two newest and most differentiated projects discussed prominently in the about section - the RAG chatbot and the multi-agent diagnostic system - are not pinned to featured, so a visitor skimming the top of the profile misses the work the about section leads with.
- **RECRUITER** — Two featured GitHub projects reinforce the technical positioning, but neither has engagement or date metadata available (consistent with LinkedIn typically not exposing analytics for featured items, and missing_fields is empty, so this is not a scraper gap).

Projects

- **CAREER COACH** — The 'AI & RAG Data Systems' headline claim and the about section's AI-native systems paragraph are entirely supported by personal open-source projects dated Aug-Sep 2026, not by EPAM job responsibilities. The about section is transparent about this ('Beyond production work'), but the headline alone doesn't carry that distinction.
- **RECRUITER** — 6 projects with GitHub links and quantified, detailed descriptions (e.g., '80 columns documented automatically,' 'processing billions of records') substantially reinforce AI/RAG/Fabric/Databricks keywords beyond what appears in top_skills or the headline, adding credibility for a human reviewer even though an automated ATS keyword scan may weight top_skills/headline more heavily than project prose.

Recommendations

- **CAREER COACH** — All three received recommendations describe the Blackboard analyst era (2019-2024); none reference the current EPAM Senior Data Engineer role, Fabric architecture work, or the data governance framework the about section leads with. External validation currently lags a full role behind the profile's stated positioning.
- **RECRUITER** — All 3 recommendations are specific and substantive rather than generic (they cite concrete deliverables like the GTM semantic model and proactive insight-finding), but all three are from the Blackboard era; none cover the current, most senior role (Senior Data Engineer at EPAM, held since Apr 2024), leaving the highest-stakes claims in the profile without third-party corroboration.
- **RECRUITER** — Jason Bass's recommendation states Felipe worked 'at Blackboard and then Anthology,' but the experience section lists both Blackboard-era roles (Nov 2019 - Mar 2024) solely under 'Blackboard,' with no mention of the Anthology rebrand. This mismatch could confuse a recruiter cross-checking employment history or company names.

Posts & Visibility

- **CAREER COACH** — Posts consistently reinforce a hands-on technical-practitioner brand (skew testing, memory ceilings, execution plans). Engagement is uneven: two fundamentals-explainer posts (1mo) reached 44K-48K impressions while the five most recent posts, including newer RAG/agent content, reach only 400-2,300 impressions - a real signal worth noting even though it sits outside pure narrative critique.

- **RECRUITER** — 7 posts spanning roughly the last month at a near-weekly cadence (2d, 1w, 2w, 3w, 4w, 1mo, 1mo) signal an actively engaged profile rather than a stale one.
- **RECRUITER** — All post topics (PySpark data skew, RAG over a Databricks lakehouse, Spark memory limits, Fabric Eventstream vs. KQL, Fabric Warehouse MERGE dedupe, large-file chunking, PySpark join execution plans) are tightly aligned with the headline and top_skills, reinforcing the hands-on Data Engineer positioning rather than diluting it with off-topic content.
- **RECRUITER** — Engagement is inconsistent: five of seven posts get modest traction (3-22 likes, 0 comments, 397-2318 impressions) while two spike heavily (178 likes/47,825 impressions and 201 likes/44,208 impressions). Comments are near-zero across nearly all posts even on the high-impression ones, suggesting content is viewed but rarely discussed.

Header Photo/Banner

- **RECRUITER** — Both photo_present and banner_present are true, removing a common 6-second-skim red flag.

PRIORITIES

Prioritized Recommendations

High Impact

Company-naming inconsistency: Blackboard vs. Anthology (High Impact)

Both agents independently flag the same mismatch: Jason Bass's recommendation states Felipe worked 'at Blackboard and then Anthology,' but the experience section lists both Blackboard-era roles (Nov 2019 - Mar 2024) solely under 'Blackboard,' with no rebrand note anywhere. Career-coach frames it as an unreconciled narrative detail; recruiter frames it as a red flag for a recruiter cross-checking employment history. Flagged from two independent angles, this is prioritized as high impact.

Before

Blackboard (company name listed for both the Senior Data Analyst and Data Analyst roles, Nov 2019 - Mar 2024)

After

Blackboard (later Anthology)

Top Skills list lags behind the profile's newest AI/RAG and Databricks work (High Impact)

Career-coach flags this as a narrative gap: the headline leads with 'AI & RAG Data Systems' as a positioning pillar, but the pinned 5-slot top_skills list contains nothing AI/RAG-related. Recruiter flags the same section as a searchability gap: Databricks appears heavily in the About section, 2 of 6 projects, and several recent posts, but is absent from top_skills, so a recruiter sourcing specifically on 'Databricks' may not surface the profile. Both agents point at the same redundant slot ('Data Engineering,' which duplicates the headline job title) as the one to swap. Flagged by both agents on the same section, this is prioritized as high impact.

Before

["SQL", "Python (Programming Language)", "PySpark", "Microsoft Fabric", "Data Engineering"]

After

["SQL", "Python (Programming Language)", "PySpark", "Microsoft Fabric", "Databricks"]

Recommendations received lag a full role behind the current position (High Impact)

Both agents independently flag that all three received recommendations describe the Blackboard analyst era (2019-2024); none reference the current EPAM Senior Data Engineer role, the Fabric architecture work, or the data governance framework the About section leads with. Career-coach frames this as external validation lagging the stated positioning; recruiter frames it as the highest-stakes claims in the profile lacking third-party corroboration. Flagged by both agents on the same gap, this is prioritized as high impact.

Medium Impact

Fabric-only positioning is in tension with recent Fabric + Databricks project mix (Medium Impact)

Career-coach notes that the headline and About section position Felipe exclusively as a Microsoft Fabric specialist, yet the two most recent flagship projects (Aug-Sep 2026) are built on Databricks rather than Fabric, creating tension between the stated platform specialization and the newest, most visible project work.

Before

Senior Data Engineer | Microsoft Fabric · PySpark · Python · SQL | Medallion Lakehouse · Real-Time Streaming | AI & RAG Data Systems

After

Senior Data Engineer | Microsoft Fabric & Databricks · PySpark · Python · SQL | Medallion Lakehouse · Real-Time Streaming | AI & RAG Data Systems

About section needs a stronger hook and a closing call-to-action (Medium Impact)

Career-coach notes the opening line ('8+ years building the data infrastructure that powers strategic decisions') is serviceable but generic rather than leading with a concrete number, and the About section ends on certification facts with no call to action inviting connection, conversation, or a look at the featured work.

Before

Opens: "8+ years building the data infrastructure that powers strategic decisions..." / Closes: "...covering both the engineering and analytics layers of the Fabric stack."

After

Opens: "100+ curated Delta Tables, billions of records processed — I build the data infrastructure that powers strategic decisions..." / Closes: "...covering both the engineering and analytics layers of the Fabric stack. Open to conversations about Fabric-based lakehouse architecture and AI-native data systems - feel free to connect or check out the projects below."

Featured items don't reflect the newest AI/RAG projects the About section leads with (Medium Impact)

Career-coach notes the two pinned featured items are both Fabric-only medallion/streaming projects, while the two newest and most differentiated projects discussed prominently in the About section - the RAG chatbot and the multi-agent diagnostic system - are not pinned, so a visitor skimming the top of the profile misses the work the About section leads with.

Before

Featured: "Real-Time Stock Volatility Pipeline – Microsoft Fabric (GitHub)" and "NYC Taxi Medallion Architecture – Microsoft Fabric (GitHub)"

After

Featured: "Real-Time Stock Volatility Pipeline – Microsoft Fabric (GitHub)" and "Medallion Data Dictionary - RAG Chatbot: A RAG chatbot that auto-generates and answers questions about a data dictionary over a Databricks Medallion (Bronze/Silver/Gold) lakehouse - 80 columns documented automatically (GitHub)"

Non-technical degree carries some ATS technical-filter risk (Medium Impact)

Recruiter notes the bachelor's degree is in Business Administration (Business/Managerial Economics) rather than a technical field, which could cause a rigid ATS filter screening for a technical degree to deprioritize the profile. The recruiter's own suggestion is that no change to the degree itself is needed - the certifications (DP-700, DP-600, PL-300) and project GitHub links already largely offset this, and should simply stay prominent.

Low Impact

Recent posts get modest comment engagement despite good reach (Low Impact)

Recruiter notes engagement is inconsistent: five of seven posts get modest traction while two spike heavily (178 and 201 likes), but comments are near-zero across nearly all posts even on the high-impression ones, suggesting content is viewed but rarely discussed.

STRATEGY

Career Trajectory — 3-Year Outlook

Forward-looking strategic guidance on the candidate's career trajectory — distinct from the profile-polish recommendations above.

SWOT Analysis

Strengths

- Holds DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate), both current through 2027, covering both the engineering and analytics layers of one platform.
- Every role shows quantified, outcome-based achievement framing: 35% latency reduction, 55% defect-rate cut, and $45M-$10M in attributed business revenue.
- Recent projects show a deliberate build-up: Fabric medallion depth through early 2026, then agentic AI systems with Claude Code, MCP, and RAG by fall 2026.
- Already leads a 3-person data engineering team at EPAM, giving early people-management exposure alongside deep individual-contributor technical work.

Opportunities

- AI Engineer roles show a 3.2-to-1 demand-to-candidate gap, and LLM-focused engineers earn 25-40% more than generalist ML engineers in 2026.
- Hiring for AI integration roles grew 178% year over year, and MCP - a skill Felipe's projects already use - is now the enterprise production standard.
- Enterprises increasingly run hybrid Fabric-plus-Databricks stacks (Mastercard, PepsiCo), rewarding engineers credible across both rather than pure single-platform specialists.

Weaknesses

- Skills list includes Databricks, Azure, and Apache Spark, but no certification backs any of them - only the Microsoft Fabric family is certified.
- The RAG chatbot and multi-agent diagnostic system remain personal open-source projects, not job-scoped production work - no MLOps or deployment experience is listed.
- Projects show MCP and basic RAG, but no listed exposure to agent orchestration frameworks like LangGraph, evaluation design, or guardrail tooling.
- Team-lead experience is a single data point - 3 engineers, current role only - with no broader engineering-management track record yet.

Threats

- Microsoft Fabric holds only about 12% platform market share in a fragmented five-platform field, so pure-Fabric specialization narrows the addressable job market.
- LangChain now appears on so many resumes it's 'no longer a differentiator,' so Felipe's RAG/agent project work alone may not stand out.
- Analysts forecast 60% of AI projects failing in 2026 on poor data readiness, which could push employer budgets toward foundational data work over exploratory RAG projects.

Recommended Paths

Stay IC and deepen into AI/Agentic Data Engineering

Rationale:

This is a judgment call among several reasonable directions, not the objectively correct one. AI Engineer demand and comp are running ahead of Data Engineer, and Felipe's RAG/multi-agent projects already sit on solid Fabric/Databricks lakehouse depth most data engineers lack.

Tradeoffs:

Requires converting GitHub project skills into actual paid job scope, and needs to move past basic LangChain/RAG into LangGraph/MCP-level depth since those are becoming table stakes rather than differentiators.

Add real Databricks production depth alongside Fabric

Rationale:

Fabric's roughly 12% share in a fragmented, multi-platform market caps pure-Fabric specialization, and enterprises are increasingly running hybrid Fabric+Databricks stacks that reward cross-platform credibility.

Tradeoffs:

Splits focus across two platforms' fast-moving roadmaps instead of going deeper on the higher-paying AI Engineer direction, risking 'competent at two' instead of 'excellent at one.'

Pursue an Analytics/Data Engineering management track

Rationale:

Felipe already leads a 3-person team at EPAM and has a cross-functional stakeholder track record (Finance, FP&A, Marketing, Sales) since the analyst era, plus governance-framework ownership that is manager-adjacent scope.

Tradeoffs:

Management typically slows hands-on technical growth just as AI Engineer IC comp pulls ahead of both Data Engineer and analytics-manager pay, risking a missed window on the current AI Engineer opportunity.

Skill & Certification Recommendations

LangGraph / agent orchestration depth (High Priority)

MCP and LangGraph listings now pay 15-25% more than LangChain-only listings, and Felipe already has MCP and multi-agent project experience to build on.

Production observability and eval design for agent systems (Medium Priority)

Eval design, guardrails, and observability are now listed among the core skills replacing outdated AI resumes, and Felipe's approval-gated pipeline already touches guardrails but lacks formal evaluation tooling.

Databricks Certified Data Engineer Associate (Medium Priority)

skills_full_list and two 2026 flagship projects already use Databricks and PySpark, but no certification backs it, unlike the two Fabric certs that already exist.

Additional Power BI/DAX certifications (Low Priority)

Already well covered by the 2021 PL-300 (valid through 2027) and DP-600; further BI-layer certification has low marginal value against the documented AI Engineer pay premium.

LangChain (as a standalone focus) (Low Priority)

LangChain listings are now common enough to no longer differentiate candidates, so time is better spent on the LangGraph/MCP depth Felipe already has a head start on.

OODA Execution Loop

Observe

Each quarter, review whether EPAM project scope has started to include any AI/RAG/agentic work formally rather than just personal GitHub projects, check activity on the RAG chatbot and self-healing pipeline repos, and re-pull current AI Engineer vs. Data Engineer demand and Fabric-vs-Databricks adoption data to see if the SWOT gaps above have moved.

→

Orient

Compare what was observed against this SWOT: has a Weakness closed (e.g. paid Databricks or MLOps work landed), has an Opportunity opened into an actual applicable role (an AI Engineer posting explicitly naming MCP or RAG), and has a Threat worsened (LangChain saturation, Fabric share trends). Re-check the SWOT findings each cycle rather than starting from scratch.

→

Decide

Pick exactly one committed action for the quarter rather than several parallel initiatives - for example, 'ship a LangGraph-based orchestration upgrade to the self-healing pipeline project' or 'apply to three AI Engineer roles that explicitly name MCP or RAG in the posting.'

→

Act

Execute the chosen action, update the about section and featured_items to reflect the newest work, and request one new recommendation each cycle from an EPAM colleague or manager until the profile has at least one recommendation validating the current Senior Data Engineer era. Grounded in profile data (skills_full_list, certifications, experience, projects) and 3 web searches on market demand, run this session.

ROADMAP

Next Steps

This Week

- Reconcile the Blackboard/Anthology company-naming mismatch between the experience section and Jason Bass's recommendation (Recommendation: 'Company-naming inconsistency: Blackboard vs. Anthology').
- Update the pinned top_skills list to close the AI/RAG and Databricks keyword gap, e.g. swapping the redundant 'Data Engineering' slot for 'Databricks' (Recommendation: 'Top Skills list lags behind the profile's newest AI/RAG and Databricks work').

This Month

- Rewrite the About section's opening hook with a concrete figure and add a closing call-to-action (Recommendation: 'About section needs a stronger hook and a closing call-to-action').
- Reframe the headline/About positioning to acknowledge the Fabric + Databricks project mix, or add a line framing the Databricks work as deliberate cross-platform range (Recommendation: 'Fabric-only positioning is in tension with recent Fabric + Databricks project mix').
- Pin the RAG chatbot and/or self-healing multi-agent pipeline project to Featured items in place of or alongside an existing Fabric-only project (Recommendation: 'Featured items don't reflect the newest AI/RAG projects the About section leads with').
- Keep certifications and project links prominent near the top of the profile to continue offsetting the non-technical degree in automated ATS screens (Recommendation: 'Non-technical degree carries some ATS technical-filter risk').

This Quarter

- Request 1-2 new recommendations from EPAM colleagues or a manager (e.g. Uladzimir Kruhlik) that speak specifically to the Fabric architecture, governance framework, and team-lead work in the current role (Recommendation: 'Recommendations received lag a full role behind the current position').
- Adopt a habit of ending posts with a direct question to the audience to convert high-impression posts into comment-driven discussion over time (Recommendation: 'Recent posts get modest comment engagement despite good reach').