LinkedIn Profile Audit Report

Senior Analytics Engineer | Microsoft Fabric & Power BI | Medallion Lakehouse · Semantic Layer · Real-Time Analytics

Generated August 31, 2026

OVERVIEW

Executive Summary

This is a strong, coherent profile: a clean four-role upward trajectory (Junior Data Analyst -> Data Analyst -> Senior Data Analyst -> Senior Analytics Engineer), a keyword-dense headline and about section built around Microsoft Fabric / Power BI / Medallion Lakehouse, consistently quantified achievements across every role ($45M opportunity generation, $210M divestiture support, 35-70% efficiency gains), a current and complementary DP-700/DP-600 Fabric certification ladder, substantive third-party recommendations, and an active, on-topic weekly posting cadence. Both the career-coach and recruiter reviews independently converge on the same two headline issues: an unreconciled Blackboard/Anthology company-naming inconsistency between a recommendation and the experience section, and two off-theme AI-tooling skill entries (Claude Code, Multi-agent Systems) that slightly dilute an otherwise razor-sharp Fabric-specialist keyword signal. Secondary polish items include a missing call-to-action at the close of the About section, no cross-platform keywords (Databricks, Snowflake, dbt, Airflow, Kafka) for recruiters sourcing beyond the Microsoft stack, a duplicated Featured item, and near-zero comment engagement on otherwise high-reach posts. Recommended path forward: deepen into applied AI / context engineering as a Fabric-native AI Engineer while retaining Fabric depth as the core differentiator.

SNAPSHOT

Profile Snapshot

- **Header:** present
- **About:** present
- **Top Skills:** 5 captured
- **Experience:** 4 captured
- **Education:** 2 captured
- **Skills Full List:** 28 captured
- **Certifications:** 3 captured
- **Featured Items:** 3 captured
- **Recent Posts:** 4 captured
- **Recommendations Received:** 3 captured
- **Recommendations Given:** 2 captured
- **Projects:** 6 captured
- **Missing:** none — all sections captured

FINDINGS

Section-by-Section Findings

About

- **CAREER COACH** — The opening line ('8+ years building the data infrastructure that powers strategic decisions') is specific enough to earn a read - it names tenure and ties directly to business impact rather than opening with generic filler like 'passionate data professional.'

- **CAREER COACH** — The about section ends on a factual sentence about certifications with no call to action - no invitation to connect, no statement of what kind of opportunity or conversation the person is open to. This is a real branding gap: readers who finish the about section aren't told what to do next.
- **CAREER COACH** — The about section is largely free of vague filler; it backs claims with concrete figures ($45M in opportunity generation, $10M in campaign-driven revenue, 100% traceability) and names the actual stakeholder teams served (Finance, FP&A, Marketing, Strategy, Sales). This specificity is a genuine strength.
- **RECRUITER** — About section reiterates headline keywords (Medallion Lakehouse, semantic layer, real-time streaming, Fabric, Power BI, Data Governance) and explicitly names both certifications by code (DP-700, DP-600) — strong for both keyword search and a quick credibility scan.

Positioning & Narrative

- **CAREER COACH** — A received recommendation from Jason Bass states Felipe 'worked on my team for several years at Blackboard and then Anthology,' but 'Anthology' (Blackboard's post-2021 rebrand/merger name) never appears anywhere in the experience section, which lists both Data Analyst roles simply as 'Blackboard' through Mar 2024. This is an unreconciled naming inconsistency between a recommendation and the experience section, likely to confuse a reader who doesn't know the corporate history.

Experience

- **CAREER COACH** — All four experience entries lead with outcome/impact framing rather than task lists: 35% latency reduction, 40% faster anomaly response, 55% defect-rate cut, $45M opportunity generation, $210M divestiture support, 38% churn reduction, 40% forecast accuracy improvement. This is a genuine strength and one of the strongest parts of the profile - it consistently answers 'so what' rather than just describing duties.
- **RECRUITER** — Every role includes heavily quantified, specific achievements ($31M and $210M divestiture support, $45M opportunity generation, $10M campaign revenue, 35-70% efficiency/latency/defect improvements, team-size and scope figures) rather than vague duties — a strong credibility signal for a skimming recruiter.
- **RECRUITER** — Title progression (Junior Data Analyst -> Data Analyst -> Senior Data Analyst -> Senior Analytics Engineer) is specific and shows a clear upward trajectory. Dates are continuous across all four roles (Dec 2017-present) with no unexplained employment gaps.
- **RECRUITER** — A recommendation from Jason Bass states Felipe 'worked on my team for several years at Blackboard and then Anthology,' but the experience section lists only 'Blackboard' for both the Data Analyst and Senior Data Analyst roles, with no 'Anthology' entry. Blackboard's education business was rebranded/merged into Anthology, so a recruiter cross-referencing the recommendation text against the experience section could read this as an inconsistent or incomplete employer history.
- **RECRUITER** — Each role uses an intro sentence plus 4-5 bulleted, metric-led lines, which is scannable rather than wall-of-text. However, the EPAM and Blackboard entries are long overall (5 dense bullets each); a skimming recruiter will likely weight only the first 1-2 bullets heavily.

Education

- **RECRUITER** — A BBA in Business/Managerial Economics plus a Cornell Data Analytics certificate program provides reasonable context for a business-to-data career path, with no gap between education end dates and the start of the professional data career.

Skills & Keywords

- **CAREER COACH** — skills_full_list is otherwise a tightly coherent Fabric/data-engineering set (Microsoft Fabric, PySpark, Delta Lake, DAX, Medallion Architecture, etc.), but it also includes 'Claude Code' and 'Multi-agent Systems,' which trace to two 2026 side projects rather than the core Fabric/analytics-engineering narrative. For a recruiter skimming for a Fabric specialist, these two entries read as off-theme and slightly dilute an otherwise sharp positioning.

- **RECRUITER** — top_skills (SQL, Python, PySpark, Microsoft Fabric, Microsoft Power BI) closely match the tools emphasized in the headline and about section, giving strong coverage of the primary terms a recruiter would search for this role.
- **RECRUITER** — Two entries ('Claude Code', 'Multi-agent Systems') are niche AI-tooling skills unrelated to the core Analytics/Data Engineer keyword cluster (Fabric, Power BI, Spark, DAX, Delta Lake, etc.) a recruiter sourcing for this role would search, slightly diluting an otherwise tight and relevant skill list.
- **RECRUITER** — No adjacent multi-platform data-engineering keywords (e.g., Databricks, Snowflake, dbt, Airflow, Kafka) appear anywhere in the profile. This is not necessarily a gap in the candidate's actual capability, but if target job postings expect cross-platform breadth beyond the Microsoft Fabric stack, this profile would not surface in searches for those specific terms.

Certifications

- **CAREER COACH** — DP-700 (Nov 2025) and DP-600 (Sep 2024) form a recent, complementary certification ladder that directly backs the headline's engineering-and-analytics dual claim. The earlier Power BI Data Analyst Associate (Sep 2021) is now over four years old but still functions as a credible foundation rather than a gap.
- **RECRUITER** — All three certifications (Fabric Data Engineer Associate DP-700, Fabric Analytics Engineer Associate DP-600, Power BI Data Analyst Associate) are current, Microsoft-issued, and directly relevant to the stated headline/target role, with no outdated or irrelevant certs diluting the list.

Featured

- **CAREER COACH** — One of the three featured items (the PySpark join/execution-plan snippet) duplicates a post that also appears verbatim in recent_posts. With only three featured slots, re-featuring an already-visible recent post instead of curating something distinct (e.g., a certification highlight, a project demo, or a different high-performing post) is a missed opportunity to showcase more of the available portfolio.
- **RECRUITER** — Featured section highlights two GitHub Fabric/Medallion projects plus the top-performing post (199 reactions), reinforcing the same technical positioning as the headline. The two GitHub items have null date and engagement fields, so their recency/reach cannot be assessed from the scraped data.

Projects

- **CAREER COACH** — The six listed projects show a deliberate skill-building arc that mirrors the certifications (NYC Taxi Medallion Architecture and the Real-Time Stock Volatility pipeline both directly extend the DP-600/DP-700 Fabric skill set), plus adjacent ML breadth (churn prediction, market basket analysis) that echoes actual paid work at Blackboard. The two most recent 2026 projects (the LinkedIn audit and blog-digest multi-agent pipelines), however, pivot into general Claude Code agent tooling that isn't tied to the Fabric/analytics-engineering narrative the rest of the profile builds - consistent with the skills-list dilution finding above.
- **RECRUITER** — Four of six listed projects are core data-engineering work well-aligned with the headline (Medallion Fabric pipelines, churn/XGBoost, market basket analysis), but two (a multi-agent LinkedIn audit pipeline and a multi-agent blog digest pipeline) are AI-agent/automation projects tangential to the stated Analytics Engineer/Fabric positioning. This isn't off-brand given the 'AI-ready' framing in the about section, but a recruiter keyword-scanning specifically for data-engineering depth may see them as a minor dilution of focus.

Recommendations

- **CAREER COACH** — Three received recommendations (a peer, and two more senior stakeholders across two companies) and two given recommendations are all specific and substantive rather than generic praise - they cite concrete deliverables (the GTM Power BI report, the $45M revenue contribution) and consistent themes (business acumen, proactive insight-finding, reliability). This is strong, credible social proof.
- **RECRUITER** — All three recommendations are specific and substantive (naming concrete deliverables like the GTM Power BI semantic model, referencing measurable business impact and working relationships) rather than generic/boilerplate praise — a strong credibility signal.

- **RECRUITER** — Both recommendations given are detailed and specific rather than generic, suggesting genuine reciprocal professional relationships. This is a minor secondary positive signal but carries less recruiter screening weight than recommendations_received.

Posts & Visibility

- **CAREER COACH** — Engagement varies sharply across the four visible posts: two narrow-decision technical breakdowns (the 1.9GB CSV chunking post, the PySpark execution-plan post) pulled 177 and 199 likes with 47.7K and 44.1K impressions respectively, while two other posts pulled single-digit likes and under 1.3K impressions. This is a real, data-backed signal about what resonates with this person's audience, not a presentation flaw to fix.
- **RECRUITER** — Four posts spanning 1d/1w/2w/3w indicate an active, roughly weekly posting cadence, which is a positive discoverability and engagement signal for a recruiter checking recent activity rather than a stale/dormant profile.
- **RECRUITER** — All four recent posts cover Fabric/PySpark/streaming/data-engineering technical topics (Eventstream vs. KQL, Warehouse MERGE dedupe, CSV chunking strategy, PySpark execution plans), directly reinforcing the headline and top_skills positioning rather than diluting it with off-topic content.
- **RECRUITER** — Engagement is uneven across posts: two older posts reach large audiences (47,689 and 44,090 impressions with 177 and 199 likes) while the two most recent posts have much smaller reach (478 and 1,292 impressions, single-digit likes). Comments are near-zero (0-2) on every post despite high impressions on the top two, suggesting content is seen and liked but rarely sparks discussion.

Header.Headline

- **RECRUITER** — Headline is specific and keyword-dense, clearly stating target role (Senior Analytics Engineer) plus core tools (Microsoft Fabric, Power BI) and specialty terms (Medallion Lakehouse, Semantic Layer, Real-Time Analytics) that recur consistently throughout about, experience, and skills sections.

Header.Photo Present / Header.Banner Present

- **RECRUITER** — Both photo and banner are present, removing a common 6-second-skim red flag (recruiters commonly bounce off profiles with no photo).

Headline

- **CAREER COACH** — The headline clearly states seniority (Senior), platform (Microsoft Fabric & Power BI), and specialization (Medallion Lakehouse, Semantic Layer, Real-Time Analytics), which gives recruiters an immediate, unambiguous read on role and domain. This is a strength, not a gap.

Missing Fields / Overall Data Completeness

- **RECRUITER** — missing_fields is empty and every schema section (header, about, skills, experience, education, certifications, featured_items, recent_posts, recommendations, projects) is populated with real data, so no fields had to be excluded from this screening review.

PRIORITIES

Prioritized Recommendations

High Impact

Reconcile the Blackboard/Anthology employer-naming inconsistency (High Impact)

Both agents independently flag the same issue from different angles: career-coach frames it as a narrative-consistency gap (a received recommendation from Jason Bass says Felipe worked at 'Blackboard and then Anthology,' but the experience section only lists 'Blackboard'), while recruiter frames it as a credibility red flag for a recruiter cross-referencing the recommendation text against the experience section. Flagged independently by both, this jointly bumps to high impact.

Before

Experience entries list the employer simply as "Blackboard" for both the Senior Data Analyst (Jan 2022 - Mar 2024) and Data Analyst (Nov 2019 - Jan 2022) roles, with no mention of Anthology anywhere in the profile.

After

Blackboard (rebranded Anthology in 2021) — applied consistently to both Blackboard-era experience entries so the company history reads consistently against the Jason Bass recommendation.

Streamline the skills list to remove off-theme AI-tooling entries (High Impact)

Career-coach flags 'Claude Code' and 'Multi-agent Systems' in skills_full_list as diluting the tight Fabric-specialist narrative (framed as a positioning gap); recruiter independently flags the same two entries as diluting recruiter keyword-search relevance (framed as a searchability gap). Because one agent frames this as a narrative issue and the other as a recruiter-search issue, it jointly bumps to high impact.

Before

skills_full_list is a single flat list mixing core Fabric/data-engineering terms (Microsoft Fabric, PySpark, Delta Lake, DAX, Medallion Architecture, Data Warehousing) with two off-theme entries: "Claude Code" and "Multi-agent Systems."

After

Remove "Claude Code" and "Multi-agent Systems" from the flat skills_full_list and instead reference them in the About section under a distinct "AI tooling / side projects" mention, keeping the top-level skills list focused on the core Fabric/analytics-engineering keyword cluster while still surfacing the AI-adjacent work for roles where it's relevant.

Medium Impact

Add a closing call-to-action to the About section (Medium Impact)

Career-coach flags that the About section ends on a factual certifications sentence with no invitation to connect or statement of what opportunities the person is open to, leaving readers without a clear next action.

Before

"I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack."

After

"I hold both the DP-700 (Fabric Data Engineer Associate) and DP-600 (Fabric Analytics Engineer Associate) certifications, covering both the engineering and analytics layers of the Fabric stack. I'm always glad to connect with fellow data/analytics engineers or talk through Fabric architecture challenges - feel free to reach out."

Evaluate adding cross-platform keyword coverage (Medium Impact)

Recruiter flags that no adjacent multi-platform data-engineering keywords (Databricks, Snowflake, dbt, Airflow, Kafka) appear anywhere in the profile, so it would not surface in recruiter searches for those specific terms even though this may not reflect an actual capability gap.

Low Impact

Lead each experience role with its strongest quantified bullet (Low Impact)

Recruiter notes the EPAM and Blackboard entries are long (5 dense bullets each) and a skimming recruiter will weight only the first 1-2 bullets heavily, so ordering matters even though the content itself is strong.

Curate distinct Featured content and surface recency signals (Low Impact)

Career-coach notes one of three Featured items duplicates a post already visible in recent_posts, a missed opportunity given only three slots. Recruiter separately notes the two GitHub Featured items have null date/engagement fields, so recruiters can't gauge how current that work is.

Optimize posting strategy for engagement (Low Impact)

Career-coach notes the two highest-performing posts (177 and 199 likes) are concrete technical-decision breakdowns versus single-digit likes on lower-context posts, and recommends leaning further into that format. Recruiter separately notes comments are near-zero (0-2) despite high impressions/likes on the top posts, and suggests ending posts with a direct question to prompt discussion.

Reorder projects to lead with core data-engineering work (Low Impact)

Recruiter notes two of six projects (the multi-agent LinkedIn audit and blog-digest pipelines) are AI-agent/automation work tangential to the Analytics Engineer/Fabric positioning; not off-brand, but a recruiter keyword-scanning for data-engineering depth may see them as a minor dilution if listed first.

STRATEGY

Career Trajectory — 3-Year Outlook

Forward-looking strategic guidance on the candidate's career trajectory — distinct from the profile-polish recommendations above.

SWOT Analysis

Strengths

- Holds DP-700 (Nov 2025) and DP-600 (Sep 2024), a current, complementary Fabric certification ladder covering both engineering and analytics layers.
- Experience shows consistent quantified business impact: $45M opportunity generation, $210M divestiture support, 55% defect-rate cut, 35% latency reduction.
- Led a team of 3 engineers at EPAM, showing early people-leadership exposure alongside deep individual-contributor technical work.
- Six self-directed Fabric and ML projects - Medallion architectures, streaming pipelines, churn and market-basket models - show deliberate, ongoing skill-building beyond the day job.

Opportunities

- AI engineer demand grew 143% year-over-year in 2026 with roughly a 3.2-to-1 gap between open roles and qualified candidates.
- Context engineering, structuring data for AI-agent consumption, is emerging as data engineering's top 2026 skill and extends his existing semantic-layer work directly.
- Microsoft Fabric passed 31,000 customers and $2B ARR, up 60% year-over-year in 2026, growing the pool of employers needing certified specialists.

Weaknesses

- No Databricks, Snowflake, or non-Microsoft cloud exposure appears anywhere in skills or experience, narrowing platform range to a single vendor ecosystem.
- No MLOps, LLM/RAG, or agentic-framework skills are listed despite two hobby projects already using informal multi-agent Claude Code patterns.
- Skills list shows no infrastructure-as-code depth (e.g., Terraform, Docker) beyond GitHub Actions and Git for deploying Fabric pipelines.
- People-leadership experience is limited to one line ('led a team of 3') with no management title or scope progression shown.

Threats

- Databricks still leads Fabric for large-scale Spark engineering and multi-cloud MLOps work, so a Fabric-only profile risks exclusion from Databricks-first employers.
- Snowflake (35%) and BigQuery (28%) still out-share Fabric (12%) in the enterprise lakehouse market, narrowing the addressable pool for Fabric-only specialists.
- Analytics-manager and senior-IC total compensation bands overlap closely ($120K-$220K), so staying pure-IC offers little pay premium over building leadership scope now.

Recommended Paths

Deepen into applied AI / context engineering as a Fabric-native AI Engineer

Rationale:

This is my judgment call among several reasonable directions the evidence supports, not the only correct path. AI engineer demand grew 143% YoY with a 3.2-to-1 talent gap, and his Medallion/semantic-layer/governance work already resembles the 'context engineering' skillset now in demand; his own multi-agent Claude Code projects show genuine, self-directed interest here.

Tradeoffs:

Pulls focus from further Fabric certification depth and risks looking unfocused unless explicitly tied back to his data-engineering base; requires learning RAG, vector search, and agent frameworks not yet in his skills list.

Add Databricks depth to stay platform-agnostic

Rationale:

Fabric is growing fast (31,000+ customers, $2B ARR) but still trails Databricks on advanced MLOps/multi-cloud work, and both trail Snowflake and BigQuery in overall lakehouse market share; adding Databricks would widen his addressable employer pool.

Tradeoffs:

His DP-600/DP-700 investment is fresh (Sep 2024, Nov 2025) and building credible dual-platform depth takes real time, which could dilute the 'Fabric specialist' narrative that currently differentiates him.

Pursue an Analytics/Data Engineering Manager track

Rationale:

He already has a 'led a team of 3' data point plus recommendation letters praising mentorship (Diego Sanabria, Uladzimir Kruhlik), and manager-vs-IC total comp bands overlap closely, so this leverages existing signal without a large skill-acquisition detour.

Tradeoffs:

Slows the AI-engineering specialization above and pulls him away from the deep hands-on technical work his recommendations and high-engagement posts show he's strongest at and enjoys.

Skill & Certification Recommendations

Context engineering / RAG patterns for agentic data systems (High Priority)

Data engineers are increasingly expected to structure data for AI-agent consumption, and this builds directly on his existing Medallion and semantic-layer work.

Vector search and embeddings fundamentals (e.g., Azure AI Search) (Medium Priority)

RAG is now table-stakes for applied AI-engineer roles and extends naturally from his existing Azure/Fabric base.

Databricks fundamentals or an associate-level certification (Medium Priority)

It hedges against Fabric's smaller enterprise lakehouse share relative to Snowflake and Databricks and broadens his platform range beyond one vendor.

LangChain or a specific agentic-framework certification (Low Priority)

Not worth prioritizing right now - his own multi-agent Claude Code projects already demonstrate practical agent-orchestration experience without needing a framework certification.

Formal people-management training or a team-lead title (Low Priority)

Only relevant if he pursues the manager path; his 'led a team of 3' line and recommendation letters show early aptitude but no formal progression yet.

OODA Execution Loop

Observe

Quarterly, track the ratio of 'Analytics Engineer' vs 'AI Engineer' vs 'Databricks Engineer' job postings in his target market, log any recruiter inbound message themes, and review LinkedIn post engagement - his two highest performers (177 and 199 likes) were both concrete technical-decision breakdowns versus single-digit likes on lower-context posts.

→

Orient

Compare any newly acquired signal (a RAG demo shipped, a Databricks module completed, a team-lead scope expansion) against the SWOT weaknesses above to see which gap is closing fastest, and re-check the SWOT findings each quarter rather than assuming last quarter's read still holds.

→

Decide

Commit to one of the three recommended paths as the lead focus for the next quarter instead of splitting effort three ways; prioritize the AI-engineering path if recruiter/job-posting signal skews toward RAG or agent skills, or the Databricks path if postings increasingly require multi-cloud/Databricks experience.

→

Act

Ship one portfolio project per quarter that demonstrates the chosen direction (e.g., a RAG-over-Fabric-Lakeho use demo, or a Databricks migration of the existing NYC Taxi Medallion project), publish it in the higher-engagement 'technical breakdown' post style, and update the about section's CTA and skills list to reflect the chosen direction rather than leaving all paths open indefinitely.

ROADMAP

Next Steps

This Week

- Add a parenthetical noting the Blackboard-to-Anthology rebrand to both Blackboard-era experience entries to fix the naming inconsistency flagged by both agents (Reconcile the Blackboard/Anthology employer-naming inconsistency).
- Append a direct call-to-action sentence to the end of the About section (Add a closing call-to-action to the About section).

- Remove 'Claude Code' and 'Multi-agent Systems' from the flat skills_full_list and, if desired, reference them separately in the About section instead (Streamline the skills list to remove off-theme AI-tooling entries).
- Reorder the bullet points within the EPAM and Blackboard experience entries so the single strongest quantified achievement leads each role (Lead each experience role with its strongest quantified bullet).

This Month

- Swap the duplicated Featured item for distinct content (e.g., the highest-engagement post or a project screenshot) and, if possible, add a date/recency indicator to the two GitHub Featured items (Curate distinct Featured content and surface recency signals).
- Reorder the Projects section so the four core Fabric/data-engineering projects lead ahead of the two AI-agent tooling projects (Reorder projects to lead with core data-engineering work).
- Decide whether genuine experience exists with cross-platform tools (Databricks, Snowflake, dbt, Airflow, Kafka) and, if so, add them to skills/about; otherwise consciously keep the Fabric-specialist positioning as-is (Evaluate adding cross-platform keyword coverage).
- Start closing new posts with a direct question to invite comments, and continue publishing in the concrete technical-decision-breakdown format that already outperforms other post types (Optimize posting strategy for engagement).

This Quarter

- Track comment and engagement trends over the quarter as the revised posting approach (technical breakdowns plus closing questions) accumulates enough data to confirm the effect (Optimize posting strategy for engagement).
- If the medium-term review concludes cross-platform breadth is worth pursuing, accumulate genuine evidence (e.g., a Databricks module or migrated project) before adding those keywords to the profile, rather than adding unverified terms (Evaluate adding cross-platform keyword coverage).